

Tutorial Business Analytics

Homework 12

Exercise 12.3 Backpropagation II

Continuing with exercise 12.2, consider the following feed-forward neural network that consists of

- An input layer ($l = 0$) representing two-dimensional points

$$a^{[0]} = \begin{pmatrix} a_1^{[0]} \\ a_2^{[0]} \end{pmatrix}^T \in \mathbb{R}^2$$

- A hidden layer $l = 1$ with 2 hidden nodes and sigmoid activation function $g^{[1]}$
- An output layer $l = 2$ with one node and sigmoid activation function $g^{[2]}$.

- a) Given the following dataset of $n = 4$ observations (in columns) and initial parameters, perform one forward pass (calculate the *empirical risk* \mathcal{L}).

$$X = \begin{pmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix}, \quad Y = (0 \quad 1 \quad 1 \quad 1)$$

$$W^{[1]} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad b^{[1]} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \quad W^{[2]} = (1 \quad -1), \quad b^{[2]} = (0)$$

- b) Perform a partial backward pass and update the parameter $W^{[2]}$ using a gradient descent with learning rate $\alpha = 1$. Perform another forward pass on the full data. Did the risk decrease?

Exercise 12.4 [Programming Exercise]

Perform one backward pass through the network of the previous exercise and update the parameters using gradient descent. Use learning rate $\alpha = 1$.

Hint: For n observations, the full parameter derivatives of the risk function are given by:

$$\begin{aligned} dW^{[1]} &= \frac{1}{n} dZ^{[1]} X^T, & db^1 &= \frac{1}{n} \text{rowSums}(dZ^{[1]}), \\ dW^{[2]} &= \frac{1}{n} dZ^{[2]} A^{[1]T}, & db^2 &= \frac{1}{n} \text{rowSums}(dZ^{[2]}) \end{aligned}$$

with

$$dZ^{[1]} = W^{[2]T} dZ^{[2]} * A^{[1]} * (1 - A^{[1]}), \quad dZ^{[2]} = A^{[2]} - Y$$

where $*$ is element-wise multiplication and for a given expression F we use the shorthand notation $dF = \partial \mathcal{L} / \partial F$.

Tutorial Business Analytics

Homework 12

Exercise 12.3 Backpropagation II

Consider the following feed-forward neural network that consists of

- An input layer ($l = 0$) representing two-dimensional points

$$a^{[0]} = (a_1^{[0]}, a_2^{[0]})^T \in \mathbb{R}^2$$

- A hidden layer $l = 1$ with 2 hidden nodes and sigmoid activation function $g^{[1]}$
- An output layer $l = 2$ with one node and sigmoid activation function $g^{[2]}$.

- a) Given the following dataset of $n = 4$ observations (in columns) and initial parameters, perform one forward pass (calculate the *empirical risk* \mathcal{L}).

$$X = \begin{pmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix}, \quad Y = (0 \quad 1 \quad 1 \quad 1)$$

$$W^{[1]} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad b^{[1]} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \quad W^{[2]} = (1 \quad -1), \quad b^{[2]} = (0)$$

To calculate the forward pass, we apply the four formulas for $z^{[1]}, a^{[1]}, z^{[2]}, a^{[2]}$ found in (a). In fact, when applying this formula to *multiple observations at once*, we can write observations in columns as in the data matrix X and still apply the same equations. Instead of vectors $z^{[1]}, a^{[1]}, z^{[2]}, a^{[2]}$ we will then get matrices $Z^{[1]}, A^{[1]}, Z^{[2]}, A^{[2]}$ that have the individual results for each observation in their columns:

$$Z^{[1]} = W^{[1]}X + b^{[1]} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \end{pmatrix} = \begin{pmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix}.$$

(Here we've used the following slight abuse of notation: when adding a vector to a matrix (with matching number of rows), we mean adding it to each column of that matrix separately.)

We continue by applying the activation function to each element of $Z^{[1]}$:

$$\begin{aligned} A^{[1]} = \sigma(Z^{[1]}) &= \sigma \begin{pmatrix} 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} \frac{1}{1+e^{-0}} & \frac{1}{1+e^{-0}} & \frac{1}{1+e^{-1}} & \frac{1}{1+e^{-1}} \\ \frac{1}{1+e^{-0}} & \frac{1}{1+e^{-1}} & \frac{1}{1+e^{-0}} & \frac{1}{1+e^{-1}} \end{pmatrix} \\ &\approx \begin{pmatrix} 0.5 & 0.5 & 0.731 & 0.731 \\ 0.5 & 0.731 & 0.5 & 0.731 \end{pmatrix}. \end{aligned}$$

Doing the same for the second layer, we get:

$$\begin{aligned}
Z^{[2]} &= W^{[2]}A^{[1]} + b^{[2]} = \begin{pmatrix} 1 & -1 \end{pmatrix} \begin{pmatrix} 0.5 & 0.5 & 0.731 & 0.731 \\ 0.5 & 0.731 & 0.5 & 0.731 \end{pmatrix} + 0 \\
&= \begin{pmatrix} 0 & -0.231 & 0.231 & 0 \end{pmatrix},
\end{aligned}$$

$$A^{[2]} = \sigma((0 \quad -0.231 \quad 0.231 \quad 0)) \approx (0.5 \quad 0.442 \quad 0.558 \quad 0.5).$$

Using the loss function given in exercise b), we can calculate the loss of the network output $A^{[2]}$ element wise:

$$\begin{aligned}
\ell(Y, A^{[2]}) &\approx -[Y * \ln A^{[2]} + (1 - Y) * \ln(1 - A^{[2]})] \\
&= -[\begin{pmatrix} 0 & 1 & 1 & 1 \end{pmatrix} * \ln((0.5 \quad 0.442 \quad 0.558 \quad 0.5)) \\
&\quad + \begin{pmatrix} 1 & 0 & 0 & 0 \end{pmatrix} * \ln((0.5 \quad 0.558 \quad 0.442 \quad 0.5))] \\
&= -[\begin{pmatrix} 0 & 1 & 1 & 1 \end{pmatrix} * (-0.693 \quad -0.815 \quad -0.584 \quad -0.693) \\
&\quad + \begin{pmatrix} 1 & 0 & 0 & 0 \end{pmatrix} * (-0.693 \quad -0.584 \quad -0.815 \quad -0.693)] \\
&= -[(0 \quad -0.815 \quad -0.584 \quad -0.693) + (-0.693 \quad 0 \quad 0 \quad 0)] \\
&= (0.693 \quad 0.815 \quad 0.584 \quad 0.693).
\end{aligned}$$

The *risk* is given by the average loss in each observation:

$$\mathcal{L}(Y, A^{[2]}) = \frac{1}{n} \sum_{i=1}^n \ell(y_{(i)}, a_{(i)}^{[2]}) = \frac{1}{4} \cdot (0.693 + 0.815 + 0.584 + 0.693) = 0.696$$

- b) Perform a partial backward pass and update the parameter W^2 using a gradient descent with learning rate $\alpha = 1$. Perform another forward pass on the full data. Did the risk decrease?

We can perform the update step using the gradient found in exercise b).

$$dW^{[2]} = dz^{[2]} \cdot \left(\frac{\partial z^{[2]}}{\partial w_1^{[2]}} \quad \frac{\partial z^{[2]}}{\partial w_2^{[2]}} \right) = (a^{[2]} - y) \cdot a^{[1]T}$$

In exercise b), however, we only considered the backward pass for a single observation (x, y) . When dealing with multiple observations at once, we need some minor adjustments:

1. Replace $dz^{[2]}$ with $dZ^{[2]} = A^{[2]} - Y$
2. Replace $a^{[1]T}$ with $A^{[1]T}$
3. We are actually interested in the derivative of the *risk*, rather than the univariate *loss* for each observation. The *risk* is $\frac{1}{n} \sum (\text{losses})$, so the calculations above would give us the sum of gradients for individual observations. However, we're interested in the averages, so we have to add a factor of $1/n$.

In total we get:

$$dW^{[2]} = \frac{\partial \mathcal{L}}{\partial W^{[2]}} = \frac{1}{n} (A^{[2]} - Y) A^{[1]T}$$

$$\begin{aligned}
&= \frac{1}{4} \left(\begin{pmatrix} 0.5 & 0.442 & 0.558 & 0.5 \end{pmatrix} - \begin{pmatrix} 0 & 1 & 1 & 1 \end{pmatrix} \right) \begin{pmatrix} 0.5 & 0.5 \\ 0.5 & 0.731 \\ 0.731 & 0.5 \\ 0.731 & 0.731 \end{pmatrix} \\
&= \frac{1}{4} \left(\begin{pmatrix} 0.5 & -0.558 & -0.442 & -0.5 \end{pmatrix} \right) \begin{pmatrix} 0.5 & 0.5 \\ 0.5 & 0.731 \\ 0.731 & 0.5 \\ 0.731 & 0.731 \end{pmatrix} \\
&= \frac{1}{4} \begin{pmatrix} -0.718 & -0.744 \end{pmatrix} = \begin{pmatrix} -0.179 & -0.186 \end{pmatrix}
\end{aligned}$$

We can thus perform one gradient update step:

$$W_{new}^{[2]} = W^{[2]} - \alpha \cdot dW^{[2]} = \begin{pmatrix} 1 & -1 \end{pmatrix} - 1 \cdot \begin{pmatrix} -0.179 & -0.186 \end{pmatrix} = \begin{pmatrix} 1.18 & -0.814 \end{pmatrix}$$

To check whether the risk decreased, let's perform another forward pass. Everything up to $A^{[1]}$ is the same as in exercise c) because we didn't change anything in the first layer. In layer 2 we get:

$$\begin{aligned}
Z_{new}^{[2]} &= W_{new}^{[2]} A^{[1]} + b^{[2]} = \begin{pmatrix} 1.18 & -0.814 \end{pmatrix} \begin{pmatrix} 0.5 & 0.5 & 0.731 & 0.731 \\ 0.5 & 0.731 & 0.5 & 0.731 \end{pmatrix} + 0 \\
&= \begin{pmatrix} 0.183 & -0.005 & 0.455 & 0.267 \end{pmatrix}
\end{aligned}$$

$$A_{new}^{[2]} = \sigma((0.183 \quad -0.005 \quad 0.455 \quad 0.267)) \approx (0.546 \quad 0.499 \quad 0.612 \quad 0.566)$$

Resulting in losses of

$$\ell_{new} = \ell(Y, A_{new}^{[2]}) = (0.789 \quad 0.696 \quad 0.491 \quad 0.568)$$

And a new risk/average loss of

$$\mathcal{L}_{new} = \frac{1}{4} (0.789 + 0.696 + 0.491 + 0.568) = 0.635$$

As the previous risk was 0.696, we see that even updating just a single matrix of parameters improved the neural network.

Exercise 12.4 [Programming Exercise]

See R script.