

A photograph of a server room with rows of server racks filled with hardware components. The racks are filled with various electronic components, including circuit boards and modules. The room has a tiled floor and fluorescent lighting.

Distributed File Systems

Agenda

- File system basics: POSIX, ext2, etc.
- User-oriented FS: Network file systems (NFS)
- Big Data FS: GFS, (HDFS)
- Erasure coding



Distributed File Systems

File System Basics

Interaction with file systems

- POSIX – *Portable OS Interface*
 - POSIX, “*The Single UNIX Specification*”
 - Aligns with the ISO C 1999 standard (stdio.h)
 - Family of standards
- Specified by IEEE Computer Society
- Today, comprised of about 20 documents
- Abstractions for programmer to achieve platform independence (portability)
- File system interface

Basic concepts

- Files
- Directories
- Links
- Metadata
- Locks

```
chris@xr2d2 / % tree -L 1
.
├── bin
├── boot
├── cdrom
├── core
├── dev
├── etc
├── home
├── initrd.img -> boot/initrd.img-4.2.0-19-generic
├── initrd.img.old -> boot/initrd.img-4.2.0-18-generic
├── lib
├── lib64
├── lost+found
├── media
├── mnt
├── opt
├── proc
├── root
├── run
├── sbin
├── srv
├── sys
├── tmp
├── usr
├── var
├── vmlinuz -> boot/vmlinuz-4.2.0-19-generic
└── vmlinuz.old -> boot/vmlinuz-4.2.0-18-generic

21 directories, 5 files
```

File system operations

- **File** operations:
 - Open
 - Read
 - Write
 - Close
 - ...
- **Directory** operations:
 - Create file
 - Mkdir
 - Rename file
 - Rename dir
 - Delete file
 - Delete dir

POSIX Files <stdio.h>

```
FILE *fopen(const char * filename, const char * mode)
```

Modes

```
r  open text file for reading
w  truncate to zero length or create text file for writing
a  append; open or create text file for writing at end-of-file
rb open binary file for reading
wb truncate to zero length or create binary file for writing
ab append; open or create binary file for writing at end-of-file
r+ open text file for update (reading and writing)
w+ truncate to zero length or create text file for update
a+ append; open or create text file for update, writing at end-of-file

r+b or rb+ open binary file for update (reading and writing)
w+b or wb+ truncate to zero length or create binary file for update
a+b or ab+ append; open or create binary file for update, writing at end-of-file
```

```
int fflush(FILE *stream);
//Any buffered data is physically persisted
int fclose(FILE *stream);
//File flushed and closed
```

POSIX Directories <stat.h>

```
int mkdir(const char* path, mode_t mode)

/* example
mkdir("/home/aj/distributed_systems", S_IRUSR |
    S_IWUSR | S_IXUSR | S_IRWXG );

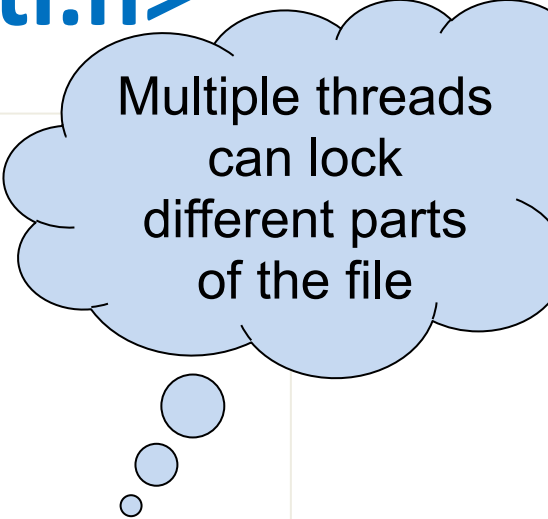
...
S_IRUSR  read permission, owner
S_IWUSR  write permission, owner
S_IXUSR  execute/search permission owner
S_IRWXG  read, write, execute/search by group
...
*/
```


POSIX File Locking <fcntl.h>

```
int fcntl(int fildes, int cmd, ...);

int fd;
struct flock fl;
fd = open("/home/aj/test.txt");
fl.l_type = F_WRLCK;           //write lock
fl.l_whence = SEEK_SET
fl.l_start = 500;              //start at byte 500
fl.l_len = 100;                //next 100 bytes

fcntl(fd, F_SETLK, &fl); //acquire lock
```



Multiple threads
can lock
different parts
of the file

Types of locks:

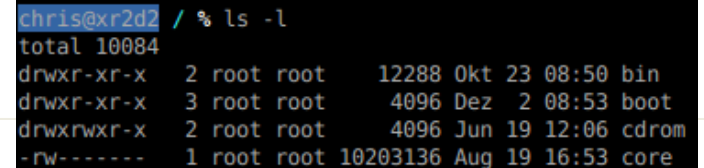
F_RDLCK	Shared or read lock
F_WRLCK	Exclusive or write lock
F_UNLCK	Unlock

POSIX File Metadata <stat.h>, <unistd.h>

```
//access permissions  
int chmod(const char * file, mode_t mode)
```

- 777 read, write, execute for all
- 664 sets **read and write** and **no execution** access for **owner** and **group**, and **read, no write, no execute** for all others

```
// set user read and write permission for file.txt  
chmod("file.txt", S_IRUSR | S_IWUSR)
```




```
chris@xr2d2 / % ls -l  
total 10084  
drwxr-xr-x  2 root root   12288 Okt 23 08:50 bin  
drwxr-xr-x  3 root root    4096 Dez  2 08:53 boot  
drwxrwxr-x  2 root root    4096 Jun 19 12:06 cdrom  
-rw-----  1 root root 10203136 Aug 19 16:53 core
```

```
//change ownership of a file  
int chown(const char *, uid_t, gid_t)
```

```
//e.g. chown("file.txt", getpwnam("arno"), -1)
```

Hard-disk Drive (HDD)

- Magnetic discs
- Cache (8MB – 128MB)
- Cost
 - ~38€/TB (1.1.2014)
 - ~30€/TB (4.12.2014)
 - ~29€/TB (7.12.2015)
 - ~22€/TB (29.11.2017)
- 5400 rpm – 15000 rpm
- Seek 4-9ms
- Connected via SATA, SCSI/SAS ...



Solid-state Drive (SSD)

- DRAM (NAND-based flash memory)
- No moving mechanical components
- Cache (16MB – 512MB)
- Cost
 - ~600€/TB (1.1.2014)
 - ~350€/TB (4.12.2014)
 - ~260€/TB (7.12.2015)
 - ~250€/TB (29.11.2017)
- Can also be connected via PCI Express
- Low-level operations differ a lot compared to HDD
 - On SSD's overwriting costs more → TRIM Command
 - Deleting is delegated to internal firmware which has a garbage collector



Source:OCZ

How common are HDD failures?


Backblaze Hard Drive Failure Rates

Ordered by Drive Size (2013 through Q3 2015)

Model Name/Number	Size	2013 Failure Rate	2014 Failure Rate	2015 Failure Rate	Failure Rate	All Periods: 2013-2015 Low Rate	High Rate
All 1.5TB Drives		16.57%	13.11%	15.10%	14.71%		
HGST(*) Deskstar 7K2000 (HDS722020ALA330)	2TB	1.03%	1.07%	2.81%	1.61%	1.40%	1.90%
Seagate Barracuda LP (ST32000542AS)	2TB	7.90%	13.43%		10.28%	6.90%	14.20%
Western Digital Red (WD20EFRX)	2TB		0.00%	6.85%	6.85%	2.40%	17.50%
All 2TB Drives		1.45%	1.42%	2.87%	1.88%		
HGST(*) Deskstar 5K3000 (HDS5C3030ALA630)	3TB	0.99%	0.59%	1.31%	0.92%	0.70%	1.10%
HGST(*) Deskstar 7K3000 (HDS723030ALA640)	3TB	1.01%	2.27%	2.12%	1.91%	1.40%	2.60%
Seagate Barracuda 7200.14 (ST3000DM001)	3TB	10.35%	43.08%	30.94%	28.46%	26.90%	29.60%
Seagate Barracuda XT (ST33000651AS)	3TB	6.91%	4.80%	3.55%	5.11%	3.50%	7.30%
Toshiba DT01ACA Series (TOSHIBA DT01ACA300)	3TB	6.93%	3.68%	2.80%	4.20%	1.40%	9.80%
Western Digital Red 3 TB (WDC WD30EFRX)	3TB	3.79%	6.94%	8.79%	7.65%	6.40%	9.30%
Western Digital Green 3 TB (WDC WD30EZRX)	3TB	6.32%	0.00%		6.32%	4.10%	9.80%
All 3TB Drives		5.22%	15.06%	4.33%	9.43%		

Hard Drive Failure Rates. All Drives. All Manufacturers.

(Year 2013-2015)



~5% fail per year

~10% failed during first 3 years

Bitrot on HDD

- Bitrot means silent corruption of data
- HDD specifications predict an **Uncorrectable bit Error Rate** (UER) of 10^{15} (1,000,000,000,000,000 ~ 125 TB)
- Evaluation [1]
 - 8x100GB HDD
 - After 2 PB reads
 - 4 read errors where observed
- How to protect against bitrot?
 - Erasure codes

[1] <http://research.microsoft.com/pubs/64599/tr-2005-166.pdf>


Disk file systems

- Linux
 - ext, ext2, ext3, ext4
 - JFS, XFS,
 - BTRFS, ZFS
 - Pooling, snapshots, checksums
- Windows
 - NTFS
 - FAT, FAT32, exFAT, ReFS
- Let's take a quick look at ext2

Linux ext2

The second extended file system


- **Superblock**, file system metadata (repeated)
 - Defines file system type, size, status, and information about other metadata structures (metadata of metadata)
- Index-nodes (**inodes**): one per file or directory
- **Data blocks**



Linux ext2

The second extended file system


- **Superblock**, file system metadata (repeated)
 - Defines file system type, size, status, and information about other metadata structures (metadata of metadata)
- Index-nodes (**inodes**): one per file or directory
- **Data blocks**



Linux ext2

The second extended file system


- **Superblock**, file system metadata (repeated)
 - Defines file system type, size, status, and information about other metadata structures (metadata of metadata)
- Index-nodes (**inodes**): one per file or directory
- **Data blocks**



Linux ext2

The second extended file system

- **Superblock**, file system metadata (repeated)
 - Defines file system type, size, status, and information about other metadata structures (metadata of metadata)
- Index-nodes (**inodes**): one per file or directory
- **Data blocks**



ext2 inode

- Owner and group identifiers
- File length
- File type and access rights
- Number of data blocks
- Array of pointers to data blocks
- Timestamp

- Types
 - File
 - Directory
 - Symbolic link

Linux/fs/ext2/ext2.h

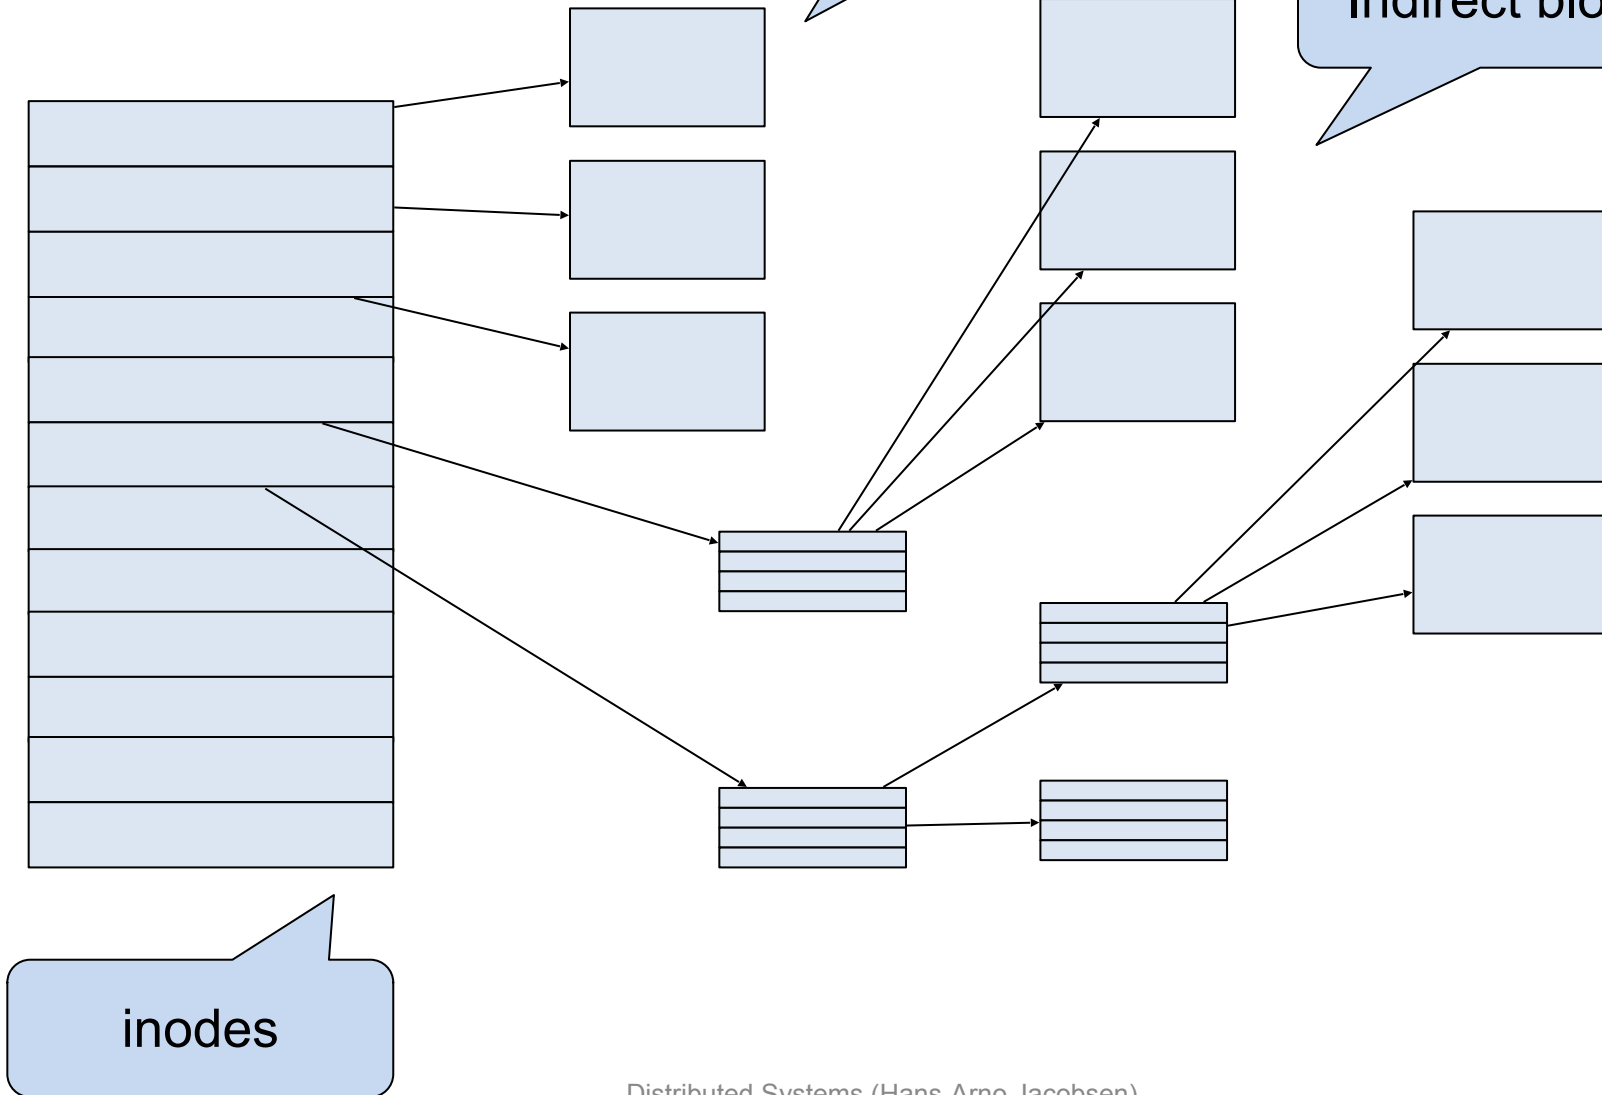
```
/*  
 * Structure of an inode on the disk  
 */  
struct ext2_inode {  
    __le16 i_mode;    /* File mode */  
    __le16 i_uid;     /* Low 16 bits of Owner Uid */  
    __le32 i_size;    /* Size in bytes */  
    __le32 i_atime;   /* Access time */  
    __le32 i_ctime;   /* Creation time */  
    __le32 i_mtime;   /* Modification time */  
    __le32 i_dtime;   /* Deletion Time */  
    __le16 i_gid;     /* Low 16 bits of Group Id */  
    __le16 i_links_count; /* Links count */  
    __le32 i_blocks;  /* Blocks count */  
    __le32 i_flags;   /* File flags */  
    ...  
  
    struct ext2_dir_entry {  
        __le32 inode;    /* Inode number */  
        __le16 rec_len;  /* Directory entry length */  
        __le16 name_len; /* Name length */  
        char name[];     /* File name, up to EXT2_NAME_LEN */  
    };  
};
```

Source: <https://github.com/torvalds/linux/blob/master/fs/ext2/ext2.h>

ext2 inode

Direct blocks

Indirect blocks





Distributed File Systems

NFS – Network File System


Distributed file systems

Motivation

- Collaboration
 - Shared file directory for projects, etc.
- Resource sharing
 - Pooling resources accross multiple devices
 - Incremental scalability (add hardware over time)
- Challenges
 - Performance
 - Scalability
 - Consistency


Distributed file system

Simplified



Distributed file system

Simplified



The Network File System (NFS)

Initially, 1984, by Sun Microsystems

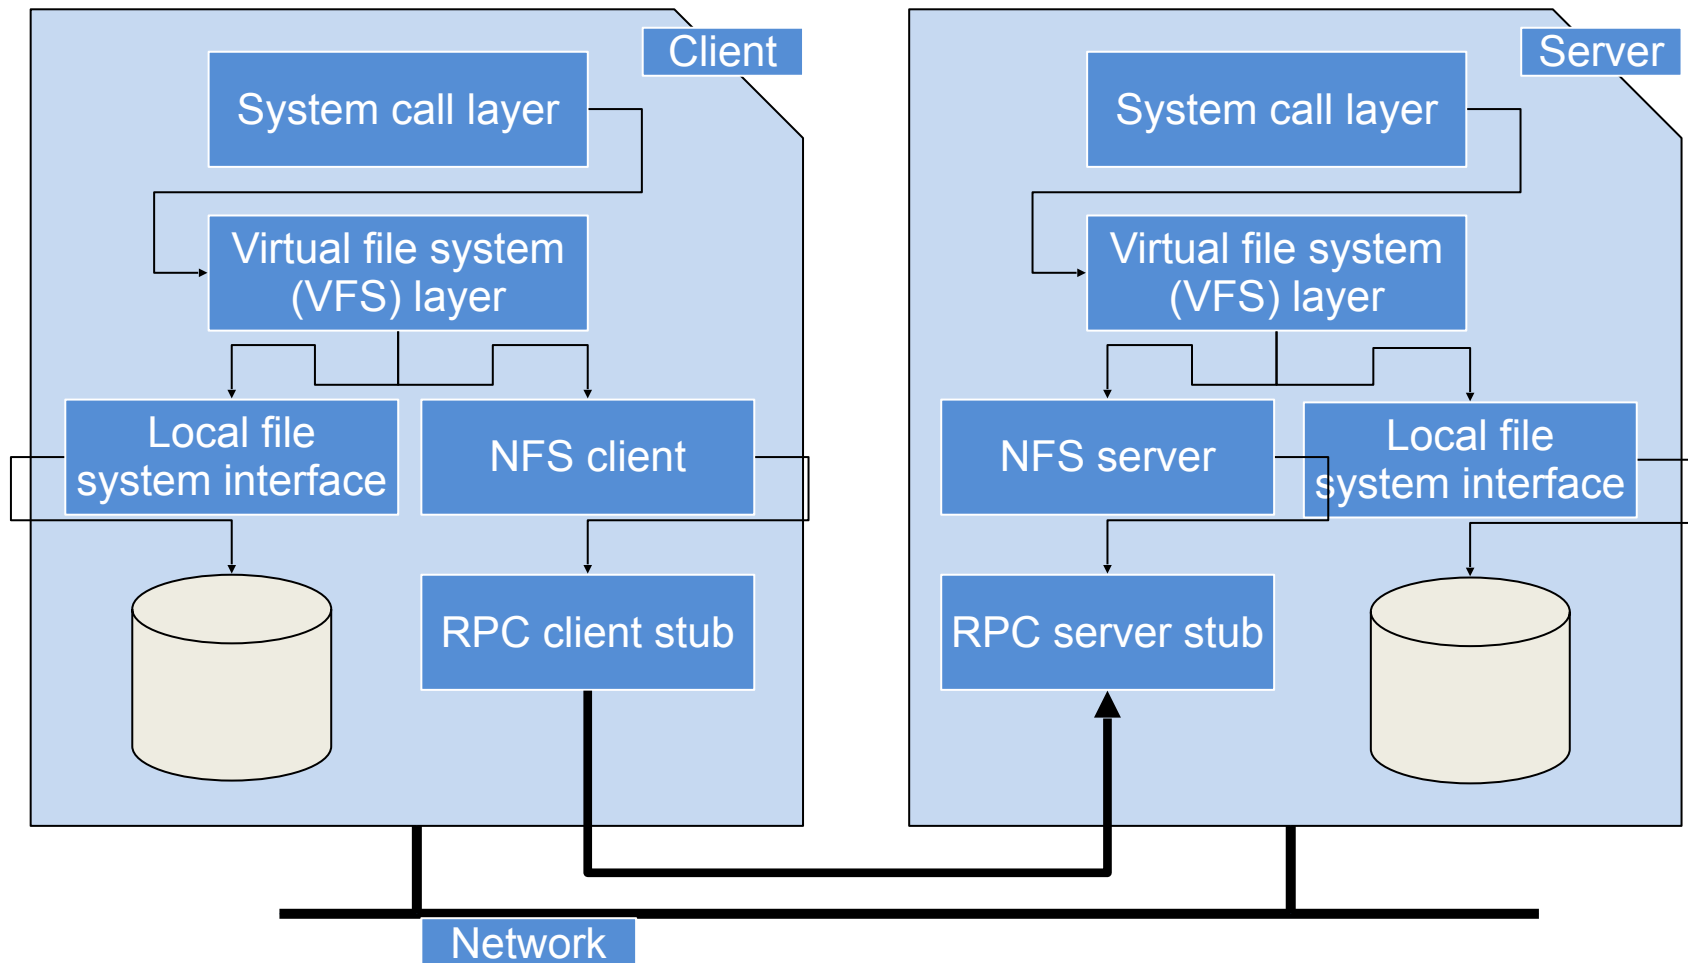
- Goals:
 - **Consistent namespace** for files across nodes
 - ***Authorized users*** can access their files from any node
- NFS protocol designed for LANs
- NFS creates a remote access layer for file systems
 - Each file is hosted on a **server** and accessed by **clients**
 - Namespace is **distributed** across servers
 - Each client treats remote files as local ones (“virtual files”)

The Network File System (NFS)

Initially, 1984, by Sun Microsystems

- NFS follows a **user-centric** design
 - Most files are privately owned by a **single user**
 - Few **concurrent access** across clients
 - Reads are **more common** than writes
- Open protocol
 - Lead to wide adoption
 - Many commercial implementation

Basic NFS architecture



Sending commands

- Essentially, NFS works as a replicated system using **remote procedure calls (RPCs)** to propagate FS operations from client(s) to server(s)
- Naïve solution: forward **every RPC** to server
 - Server orders all incoming operations, performs them, returns results
- Downside
 - High access latency due to RPCs
 - Server becomes overloaded by many RPCs

Solution: Caching

- Clients use a cache to store a copy of remote files, called “virtual files”
- Clients periodically synchronize with server
- This is essentially **multi-primary replication**:
 - *How should synchronization be done? (eager/lazy)*
 - *What is the right consistency level?*

Original version: Sun NFS

NFSv2, ..., NFSv4, ...

- Developed in 1984
- Uses in-memory caching:
 - File blocks, directory metadata
 - Stored at both clients and servers
- Advantage: no network traffic for open, read, write
- Problems: **failures** and **cache consistency**

Failures I

- **Server crash**
- Any data not persisted to disk is lost
- What if client does seek(); [***server crash***]; read()?
 - Seek sets a position offset in the opened file
 - After crash, **server forgets** offset, **read returns incorrect data**

Failures II

- **Communication omission failures**
 - Client *A* sends `delete (foo)`, server processes it
 - Server acknowledgement of delete is lost, meanwhile Client *B* issues `create (foo)`
 - Client *A* times out and send `delete (foo)` again, **deleting the file created by Client *B*!**
- **Client crash**
 - Since caching is in memory, **lose all updates by client** if not synched to server

Solution: Stateless RPC

- RPC commands are **stateless**: server does not maintain state across commands in a “session”
- *read()* is stateful (server needs to remember *seek()*) → **read(position)** is **stateless** (server has all the information needed for correct read)
- With stateless RPC, server can fail and later continue to serve commands without recovering former state


Solution: Idempotent RPC

- NFS's RPCs are designed to be **idempotent**
- Repeating a command has no side effect
- `Delete("foo")` becomes **`delete(someid)`**, so it cannot wrongly delete a new file named "foo"
- Read, lookup are idempotent
- Write includes, data to write, the file ID, the offset to write at, therefore, idempotent

Common loss scenarios

Handled by client via timeout, retry, idempotent server RPCs


Case 1: Request Lost



Case 2: Server Down



Case 3: Reply lost on way back from Server



Is mkdir idempotent?

- MKDIR message from client to server
- Server ACKing successful creation is lost
- Client times out and retries MKDIR
- Server responds with error directory exists
- NFS designers opted to keep design simple

Cache Consistency

- Clients can cache file blocks, directory metadata, *etc.*



- *What happens if both clients want to write?*

Solution: Time-bounded consistency

- Flush-on-close: When a file is **closed**, modified blocks are sent to server **synchronously** (close() does not return until update is finished)
- Each client periodically checks with server for updates
- Clients synchronize their cache **after some bounded time** if there are no more updates; otherwise they would read stale data

Concurrent Writes in NFS

- NFS does not provide any guarantees for concurrent writes!
- Server may update using one client's writes, other's writes, or a mix of both!
- Not usually a concern due to the **user-centric** design: assuming there are no concurrent writes
- A big problem if one needs to support concurrent writes

NFS Summary

- **Transparent** remote file access
- **Client-side caching** for improved performance
- **Stateless** and **idempotent** RPCs for fault-tolerance
- Periodical synchronization with the server, with **flush-on-close semantics**
- No guarantees for concurrent writes



Distributed File Systems

GFS – Google File System

The Google File System (GFS)

Design assumptions

- Designed for Big Data workloads
 - Huge files (100MB+), not optimized for small files
- Fault tolerance while running on inexpensive commodity hardware
 - 1000s machines where failure is the norm
- Introduces an API which is designed to be implemented scalably (non-POSIX)
- Architecture: one master, many chunk (data) servers; can operate across WAN links
 - Master stores metadata and monitors chunkservers
 - Chunkservers store and serve data chunks

Design assumptions

- Read workload
 - Large streaming reads (data caching not beneficial); no client-side data caches
 - Small random reads
- Write workload
 - File append via producer-consumer pattern
 - Hundreds of concurrently appending clients
 - Modification supported but not a design goal
- Bandwidth is more important than low latency


Interface

- Supported operations
 - Create, delete, open, close, read, write
 - Record append
 - Allows multiple clients to append data to the same file while guaranteeing atomicity
 - Snapshot
 - Creates a copy of a file or a directory tree at low costs
- Does not support full POSIX interface
 - POSIX requires many guarantees which are hard to fulfill in distributed applications

Architecture


- Files
 - Divided into fixed-size chunks (64MB)
 - Identified by an immutable and unique id (chunk handle)
- Single master
 - Maintains GFS metadata
 - Namespace, access control information, mapping from files to chunks, location of chunks
 - Heartbeats chunkservers
- Multiple chunkservers
 - Chunks are stored on disk as Linux files
 - Each chunk is replicated to multiple chunkservers (depending on a replication factor defaulting to 3)

GFS architecture



Metadata kept at Master

- File and chunk namespaces
- Mapping from files to chunks
- Location of each chunk's replicas



Metadata management by Master

- Replicated to a shadow master and logged to operation log
 - **Namespaces**
 - **File to chunk mapping**
- **Location of chunks** is in-memory only (fast)
 - At start-up, periodically, upon failover, master asks chunkservers which chunks they have to rebuild location-to-chunk mapping
 - Periodic scanning is used to implement
 - Garbage collection (when files are deleted)
 - Re-replication (chunkserver failure)
 - Chunk migration (to balance load and disk space)
- Metadata has to fit in memory (64 bytes/chunk)

Operation log at Master

- Maintains all file creating, renaming, deletion operations *etc.*
- Only persistent, historical record of metadata changes
- Persisted to local disk and replicated to shadow master(s)
- Metadata changes are only visible after they are persisted
- Serves for Master recovery by replaying operation log
- Periodic checkpointing of Master state to minimize replaying effort

How is fault-tolerance achieved?

- Master
 - Operation log, replication of log and metadata to shadow master
- Chunkserver
 - All chunks are versioned
 - Version number updated upon modification
 - Chunks with old versions are not served and are deleted
- Chunks
 - Re-replication triggered by master, maintains replication factor
 - Rebalancing to distribute load among chunkservers
 - Data integrity checks

How is high-availability achieved?

- Fast recovery of master
 - Checkpointing and operation log
- Shadow master(s)
 - Serve read traffic, reduces downtime during failover
- Heartbeat messages (often include piggy-backed status updates)
 - Discover chunkserver failure
 - Trigger re-replication
 - Share current load
 - Trigger garbage collection
- Diagnostic tools

Summary on GFS

- Highly concurrent reads and appends
- Highly scalable
- On cheap commodity hardware
- Built for map-reduce kind of workloads
 - Reads
 - Appends
- Developers have to understand the limitations and may have to use other mechanisms to work around
- No POSIX API, would require many guarantees which are difficult to fulfill in DS