

Peer-to-Peer Systems



Agenda

- Peer-to-peer networking
 - Definition
 - Use cases
- Peer-to-peer overlays
 - Unstructured systems
 - Distributed hash table abstraction
 - Structured systems realizing DHTs

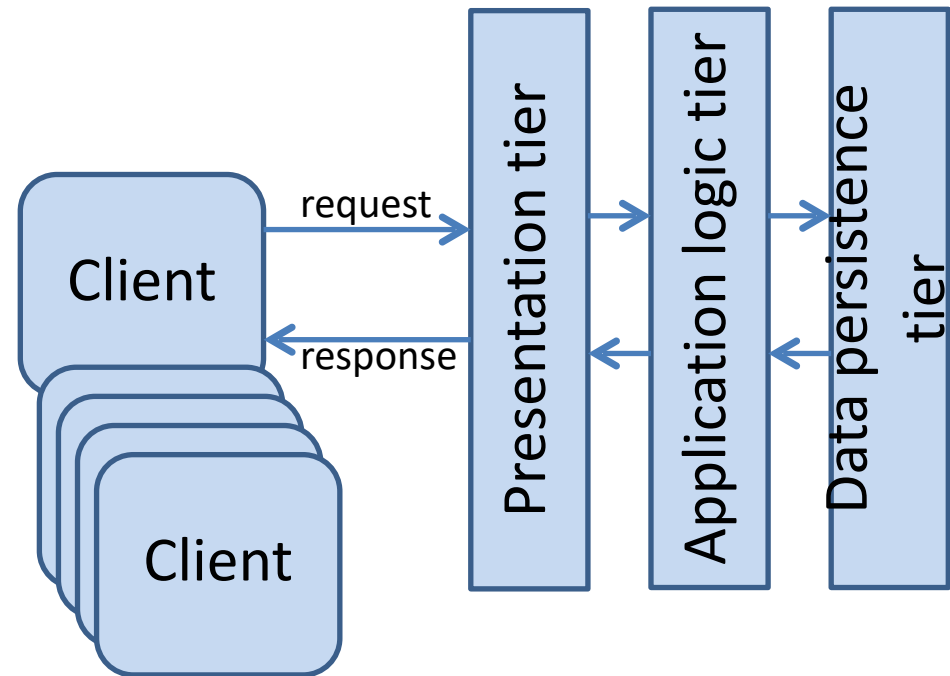

Peer-to-Peer Systems



Introduction

Client-Server Architecture

- Well known, powerful, reliable server as data source
- Clients request data from server
- Very successful model
 - WWW (HTTP), FTP, Web services, etc.
- N-tier architecture



Client-Server Limitations

- Scalability is expensive (vertical vs. horizontal)
 - Presents a single point of failure
 - Requires administration
 - Unused resources at network edge
-
- P2P systems try to address these limitations and leverage (otherwise) unused resources


Peer-to-Peer

Compute, storage, network

- P2P computing is the sharing of **compute resources** and **services** by **direct** exchange between peers (a.k.a. nodes)
- These resources and services include the **exchange of data, processing cycles, cache storage, disk storage, and bandwidth**
- P2P computing takes advantage of existing computing power, computer storage and networking connectivity, allowing users to leverage their **collective power to the ‘benefit’ of all**

* From (accessed June, 2004) http://www-sop.inria.fr/mistral/personnel/Robin.Groenevelt/Publications/Peer-to-Peer_Introduction_Feb.ppt


What is a P2P system?



- A distributed system architecture
 - No centralized control
 - Nodes are symmetric in function
- Larger number of unreliable nodes
- Enabled by technology improvements

P2P Architecture

- All nodes are both *clients* and *servers*
 - Provide and consume
 - Any node can initiate a connection
- No centralized data source
 - “*The ultimate form of democracy on the Internet*”
 - “*The ultimate threat to copyright protection on the Internet*”



- In practice, **hybrid models** are popular
 - Combination of client-server & peer-to-peer
 - Skype (early days, now unknown)
 - Spotify (peer-assisted)
 - BitTorrent (trackers)

P2P Benefits I

- Efficient use of resources
 - Unused bandwidth, storage, processing power at the edge of network
- Scalability
 - Consumers of resources also donate resources
 - Aggregate resources grow naturally with utilization
 - Organic scaling (more users, more resources)
 - “Infrastructure-less” scaling
- **Caveat:** It is not a one size fits all
 - Large companies are not switching in droves to P2P

P2P Benefits II

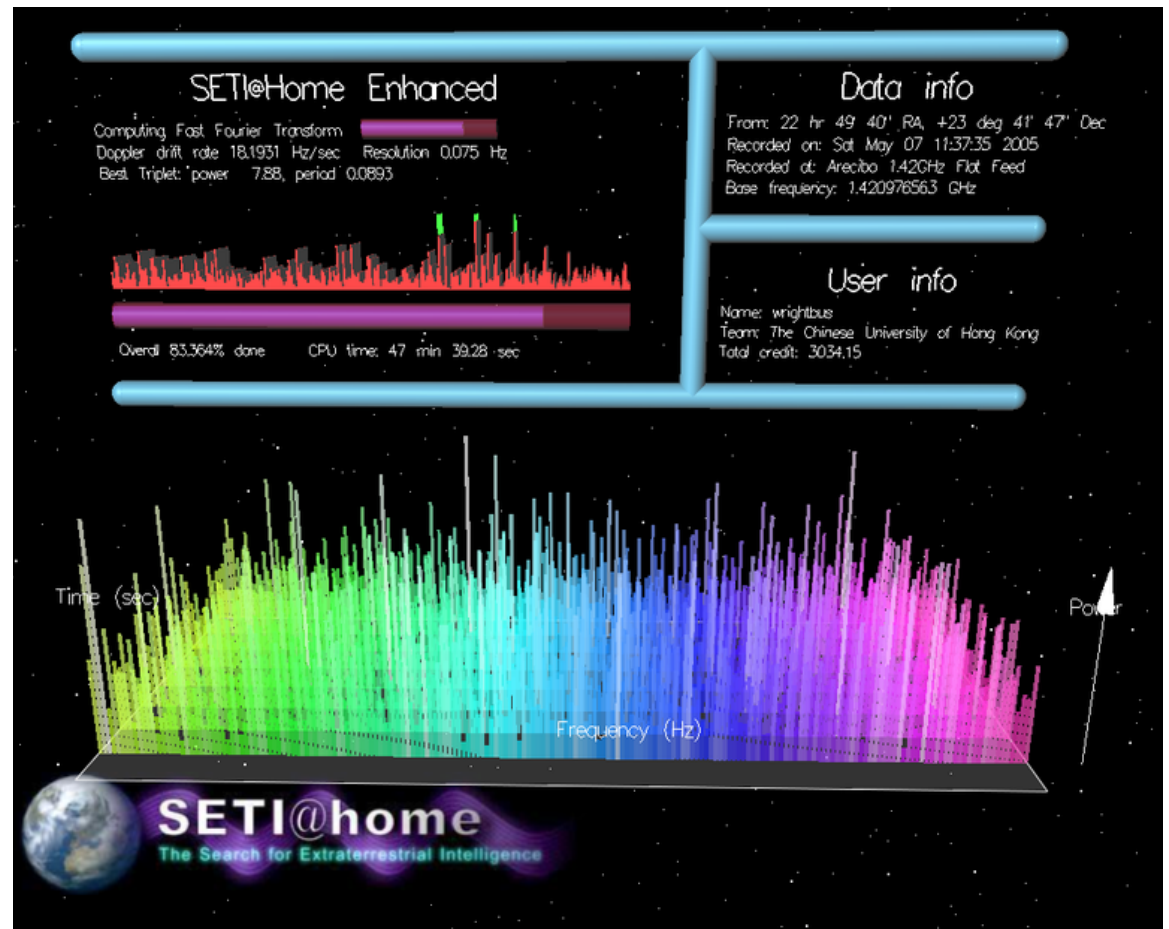
- Reliability (in aggregate)
 - Many replicas, redundancy
 - Geographic distribution
 - No single point of failure and control
- Ease of administration
 - Nodes self-organize
 - No need to deploy servers to satisfy demand
 - Built-in fault-tolerance, replication, and load balancing

Use Cases: Large-Scale Systems

- Some applications require immense resources
 - CPU: Scientific data analysis (*@home)
 - Bandwidth: Streaming, file sharing
 - Storage: Decentralized data, file sharing
- Thousands or even millions of nodes
 - How to efficiently manage such a large network?

SETI@home – started in 1999

- 5.2 million participants worldwide
- On September 26, 2001, had performed a total of 10^{21} flops
- 35 GB/data per day, 140K work units
- 30 hours/WU
- 4.2M hours of computation/day
- Centralized database



SETI@home – started in 1999


- 5.2 million participants worldwide
- On September 2001, performed 10²¹ floating point operations
- 35 GB of data downloaded to 140K computers
- 30 hours of data processed
- 4.2M hours of computation/day
- Centralized database

Start date: May 17, 1999

Project goal(s): Discover radio evidence of extraterrestrial life

Active users: 103,480 (January 2018)


Total users: 1,716,012 (January 2018)



2015 NA Traffic

Traffic share in North America during peak hours

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	28.56%	Netflix	37.05%	Netflix	34.70%
2	Netflix	6.78%	YouTube	17.85%	YouTube	16.88%
3	HTTP	5.93%	HTTP	6.06%	HTTP	6.05%
4	Google Cloud	5.30%	Amazon Video	3.11%	BitTorrent	4.35%
5	YouTube	5.21%	iTunes	2.79%	Amazon Video	2.94%
6	SSL - OTHER	5.10%	BitTorrent	2.67%	iTunes	2.62%
7	iCloud	3.08%	Hulu	2.58%	Facebook	2.51%
8	FaceTime	2.55%	Facebook	2.53%	Hulu	2.48%
9	Facebook	2.25%	MPEG - OTHER	2.30%	MPEG	2.16%
10	Dropbox	1.18%	SSL - OTHER	1.73%	SSL - OTHER	1.99%
		65.95%		78.69%		76.68%



<https://www.sandvine.com>

2015 NA Traffic

Still #1 in upstream!

Traffic share in North America during peak hours

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	28.56%	Netflix	37.05%	Netflix	34.70%
2	Netflix	6.78%	YouTube	17.85%	YouTube	16.88%
3	HTTP	5.93%	HTTP	6.06%	HTTP	6.05%
4	Google Cloud	5.30%	Amazon Video	3.11%	BitTorrent	4.35%
5	YouTube	5.21%	iTunes	2.79%	Amazon Video	2.94%
6	SSL - OTHER	5.10%	BitTorrent	2.67%	iTunes	2.62%
7	iCloud	3.08%	Hulu	2.58%	Facebook	2.51%
8	FaceTime	2.55%	Facebook	2.53%	Hulu	2.48%
9	Facebook	2.25%	MPEG - OTHER	2.30%	MPEG	2.16%
10	Dropbox	1.18%	SSL - OTHER	1.73%	SSL - OTHER	1.99%
		65.95%		78.69%		76.68%



<https://www.sandvine.com>

2015 NA Traffic

Still #1 in upstream!

Traffic share in North America during peak hours

But streaming is larger overall...

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	28.56%	Netflix	37.05%	Netflix	34.70%
2	Netflix	6.78%	YouTube	17.85%	YouTube	16.88%
3	HTTP	5.93%	HTTP	6.06%	HTTP	6.05%
4	Google Cloud	5.30%	Amazon Video	3.11%	BitTorrent	4.35%
5	YouTube	5.21%	iTunes	2.79%	Amazon Video	2.94%
6	SSL - OTHER	5.10%	BitTorrent	2.67%	iTunes	2.62%
7	iCloud	3.08%	Hulu	2.58%	Facebook	2.51%
8	FaceTime	2.55%	Facebook	2.53%	Hulu	2.48%
9	Facebook	2.25%	MPEG - OTHER	2.30%	MPEG	2.16%
10	Dropbox	1.18%	SSL - OTHER	1.73%	SSL - OTHER	1.99%
		65.95%		78.69%		76.68%



<https://www.sandvine.com>

P2P File Sharing Systems

- Large-scale sharing of files
 - User *A* makes files (music, video, etc.) on their computer available to others
 - User *B* connects to “network,” searches for files and downloads files *directly* from User *A*
- P2P networks
 - Peers are connected to each other to form an **overlay network**
 - Peers communicate using links established in overlay
- Fallen out of favor (*RIP 1999-2015*)
 - Issues of copyright infringement
 - Cloud infrastructures has taken over (controlled resources)
 - Harder to exploit mobile and connected devices
 - Streaming makes file sharing obsolete (cf. P2P Streaming)

Types of P2P Systems

- Unstructured networks
 - No obvious structure in overlay topology
 - Peers simply connect to anyone in existing network
- First generation:
 - Centralized: **Napster**
 - Pure: **Gnutella**, Freenet
- Second generation:
 - Dynamic “supernodes”
 - Hybrid: **Skype**, **Spotify**, FastTrack, eDonkey, **BitTorrent**
- Structured networks
 - Topology of overlay is controlled: peers’ connections are fixed
 - Based on the distributed hash table abstraction (**DHT**)

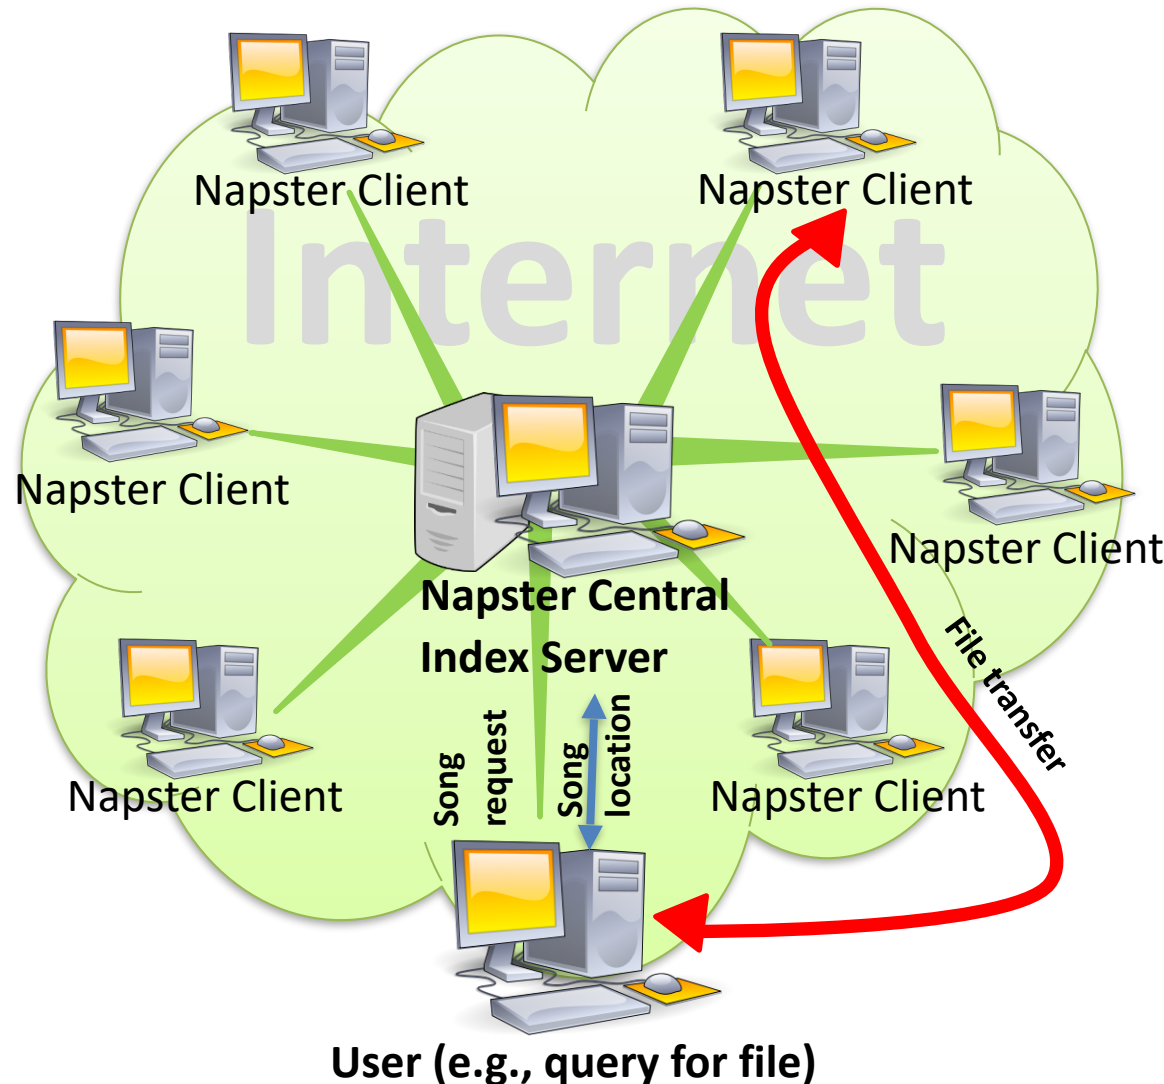
Peer-to-Peer Systems




Unstructured peer-to-peer systems

Centralized: Napster “June 1999 – July 2001”

- Centralized search indexes music files
 - Perfect knowledge
 - Bottleneck
- Users query server
 - Keyword search (artist, song, album, bit rate, etc.)
- Napster server replies with IP address of users with matching files
- Querying users connect **directly** to remote node for file to download



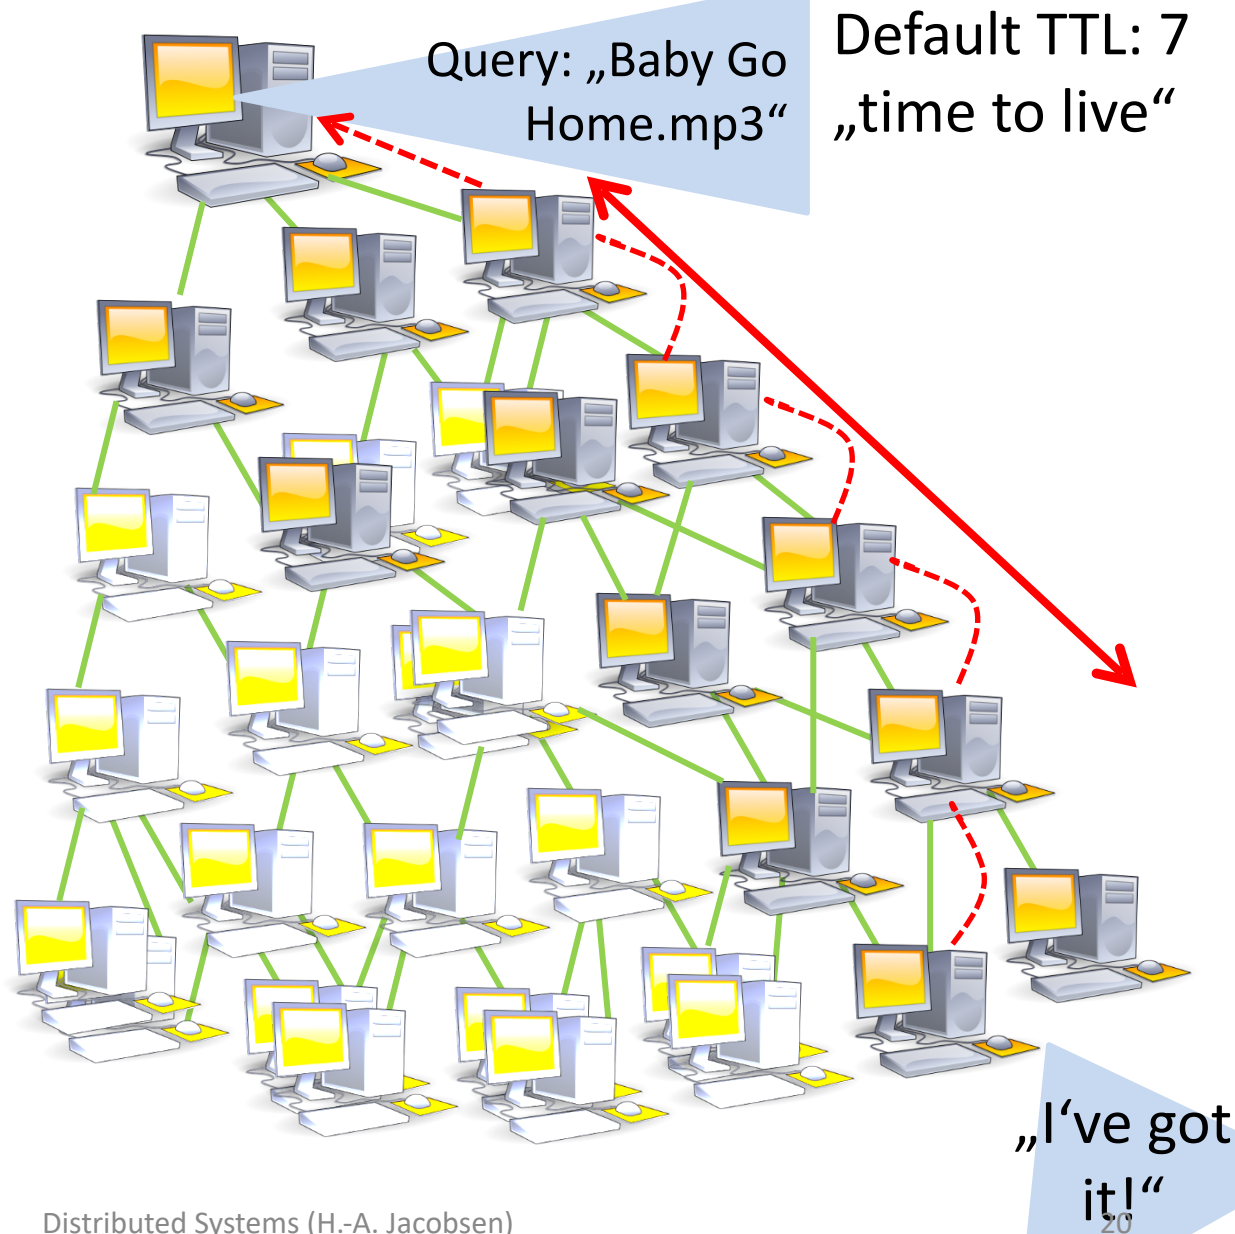
The Rise and Fall of Napster



RIAA shut down Napster easily because it used a central server!

Pure P2P: Gnutella 0.4 (2000 – 2001)

- Share any type of files
- Decentralized search
 - **Imperfect** content availability
- Client connect to (on average) 3 peers
- **Flood a QUERY** to connected peers
- Flooding propagates in network up to TTL
- Users with matching files **reply with QUERYHIT**
- Flooding wastes **bandwidth**: Later versions used more sophisticated search



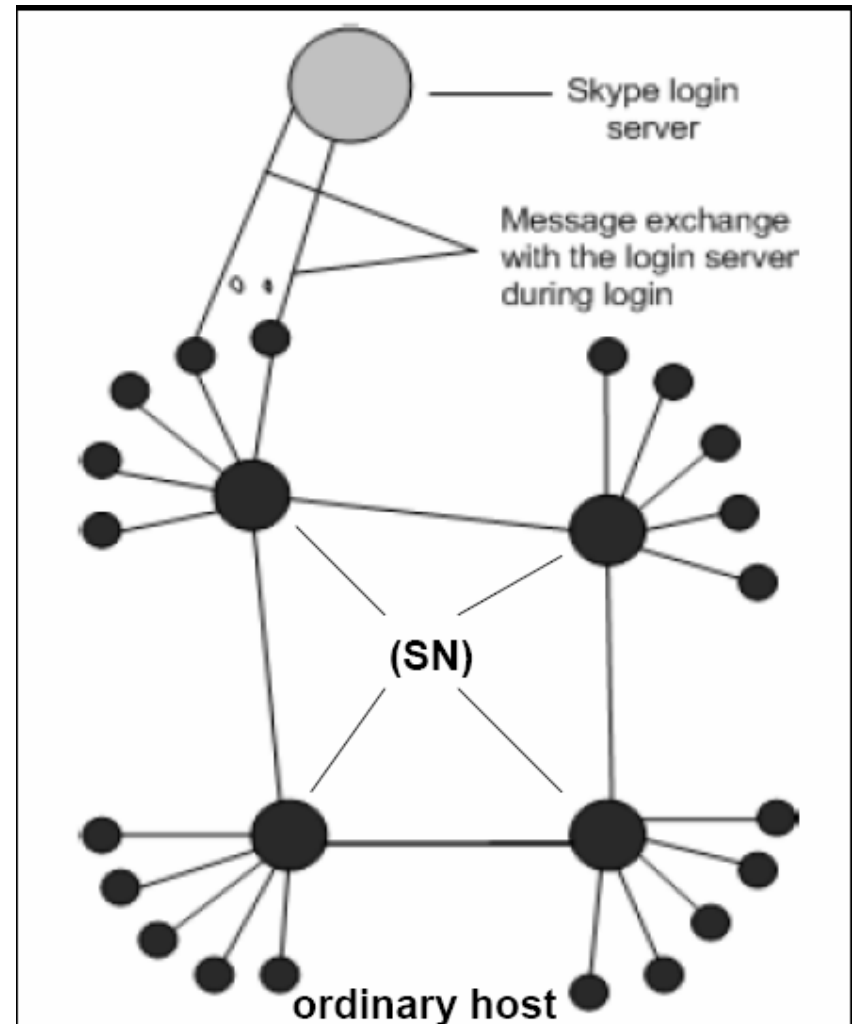
Hybrid P2P

- Both previous approaches have advantages and disadvantages
 - Centralized: single point of failure, easy to control, but perfect content availability
 - Pure: decentralized (resistant), costly and unreliable search
- Hybrid P2P combines both approaches
 - Hierarchy of peers
 - Superpeers with more capacity, discovered **dynamically**
 - Normal peers (leaf nodes) are users
- Superpeer responsibilities
 - Participates in search protocol, indexes and caches data
 - Improves content availability
 - Reduces message load

Skype Network

Around 2004

- **Super Nodes:** Any node with a public IP address having sufficient CPU, memory and network bandwidth is candidate to become a superpeer
- **Ordinary Host:** Host needs to connect to superpeer and must register itself with the Skype login server
- Login server and PSTN gateway (not shown) **are centralized components**



Responsibilities of Superpeers

In Skype Around 2004

- Indexes user directory
 - Distributed among superpeers
 - Communication among superpeers for lookup
- Communication relay
 - NAT traversal
- Phased out by Microsoft in 2011 (speculation)
 - Replaced with private servers
 - “[T]hat is in part why Skype has switched to server-based “dedicated supernodes”... nodes that **we control**, can handle orders of magnitudes **more clients per host**, are in **protected data centers** and up all the time, and **running code that is less complex** than the entire client code base. ”

Skype Impact

- Skype has shown, at least has suggested,:
 - **Signaling**, unique property of traditional phone system, is accomplished effortlessly with self-organizing P2P networks
 - **P2P overlay networks can scale** up to handle large-scale connection-oriented real-time services such as video and voice
- AT&T: *“The end of landlines ...”*



Peer-to-Peer Systems



DHT - Distributed Hash Table

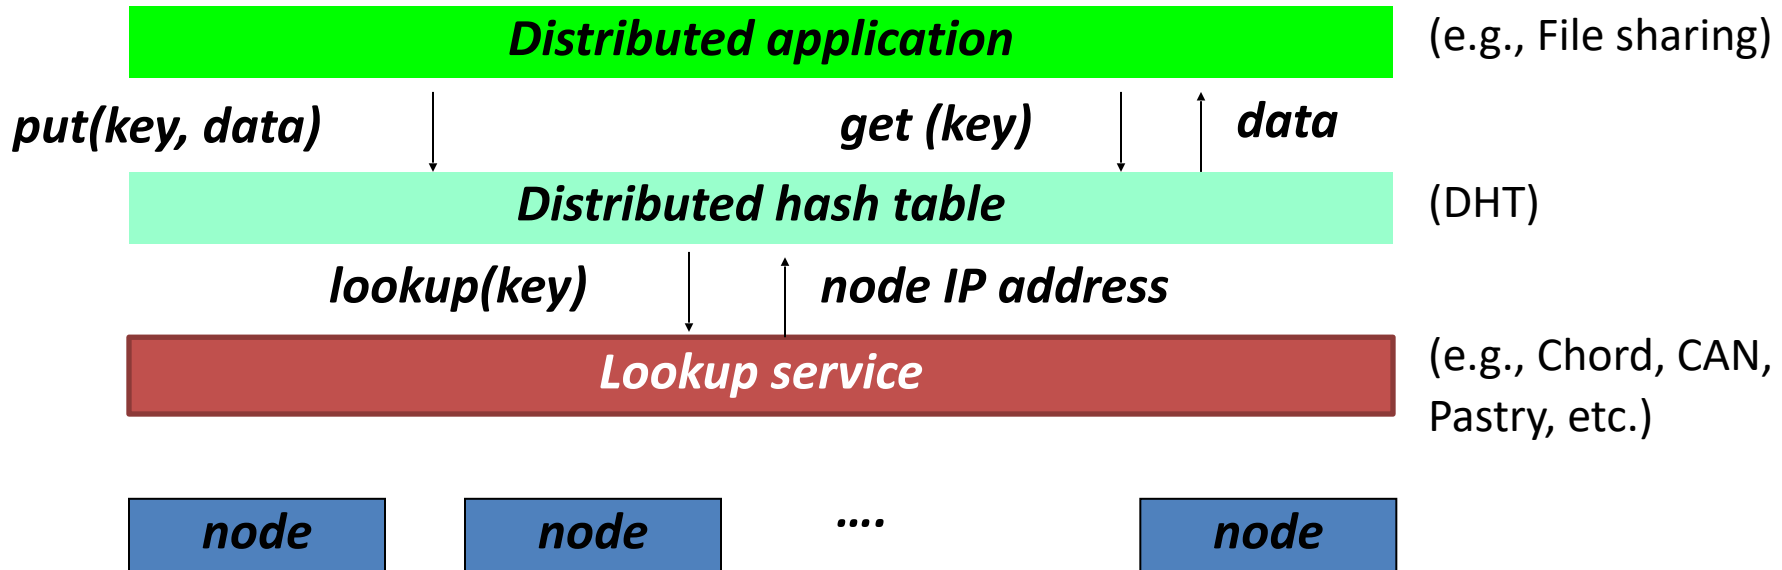
Structured Peer-to-Peer Systems

- Third generation P2P overlay networks
- Self-organizing, load balanced, fault-tolerant
- “Fast” and efficient **lookup guarantees**, e.g.,
 - $O(\log(n))$ lookups
 - $O(1)$: centralized, $O(n)$: pure, $O(\#S)$: hybrid
- Based on a hash table interface (cf. KV-Store)
 - Put(Key, Data) and Get(Key)
 - Coined term **distributed hash table** (DHT)
- Systems: Chord, CAN, Pastry, etc.

Distributed Hash Tables (DHT)

- Distributed version of a hash table data structure
- Store and retrieve (key, value)-pairs
 - **Key** is like a filename, hash of name, hash of content (since name could change)
 - **Value** is file content
 - Often just a reference to a node with the content
- Keys are hashed and mapped to a set of distributed nodes
 - Realization via consistent hashing *et al.*
 - System change should impact few nodes

DHT Abstraction



- Application distributed over many nodes
- DHT distributes data storage over many nodes

DHT Interface

- Put(key, value) and get(key) → value
 - Simple interface!
- API supports a wide range of applications
 - DHT imposes neither structure nor meaning on keys/values
- Key-value pairs are persisted and globally available
 - Good availability, content stored at edge
 - Store keys in other DHT values
 - Thus, build complex data structures

DHT as Infrastructure or Service

- Many applications can share single DHT service
- Eases deployment of new applications
- Pools resources from many participants (P2P...)
- Essentially, a middleware service, a piece of distributed systems infrastructure

DHT-based Projects

- File sharing [CFS, OceanStore, PAST, Ivy, ...]
- Web cache [Squirrel, ..]
- Archival/Backup store [HiveNet, Mojo, Pastiche]
- Censor-resistant stores [Eternity,..]
- DB query and indexing [PIER, ...]
- Event notification (Publish/Subscribe) [Scribe, ToPSS]
- Naming systems [ChordDNS, Twine, ..]
- Communication primitives [I3, ...]
- Key-value stores [Cassandra*, Dynamo*, ...]

Common denominator:

- **Data is location-independent**
- **All leverage DHT abstraction**

* In as far as they use consistent hashing among nodes

Peer-to-Peer Systems




Chord DHT

DHT Requirements

- Keys mapped evenly to all nodes in the network (**load balancing**)
- Node arrivals & departures only affect a few nodes (**low maintenance**)
- Each node maintains information about only a few other nodes (**low maintenance**)
- Messages can be routed to a node efficiently (**fast lookup**)

Chord Identifier Circle

- **Nodes** organized in an **identifier circle** based on **node identifiers**
- **Keys** assigned to their **successor** node in the identifier circle
- **Hash function** ensures **even distribution of nodes and keys** on the identifier circle
- Cf. **consistent hashing**
- With N nodes and K keys each **node is responsible for roughly K/N keys**



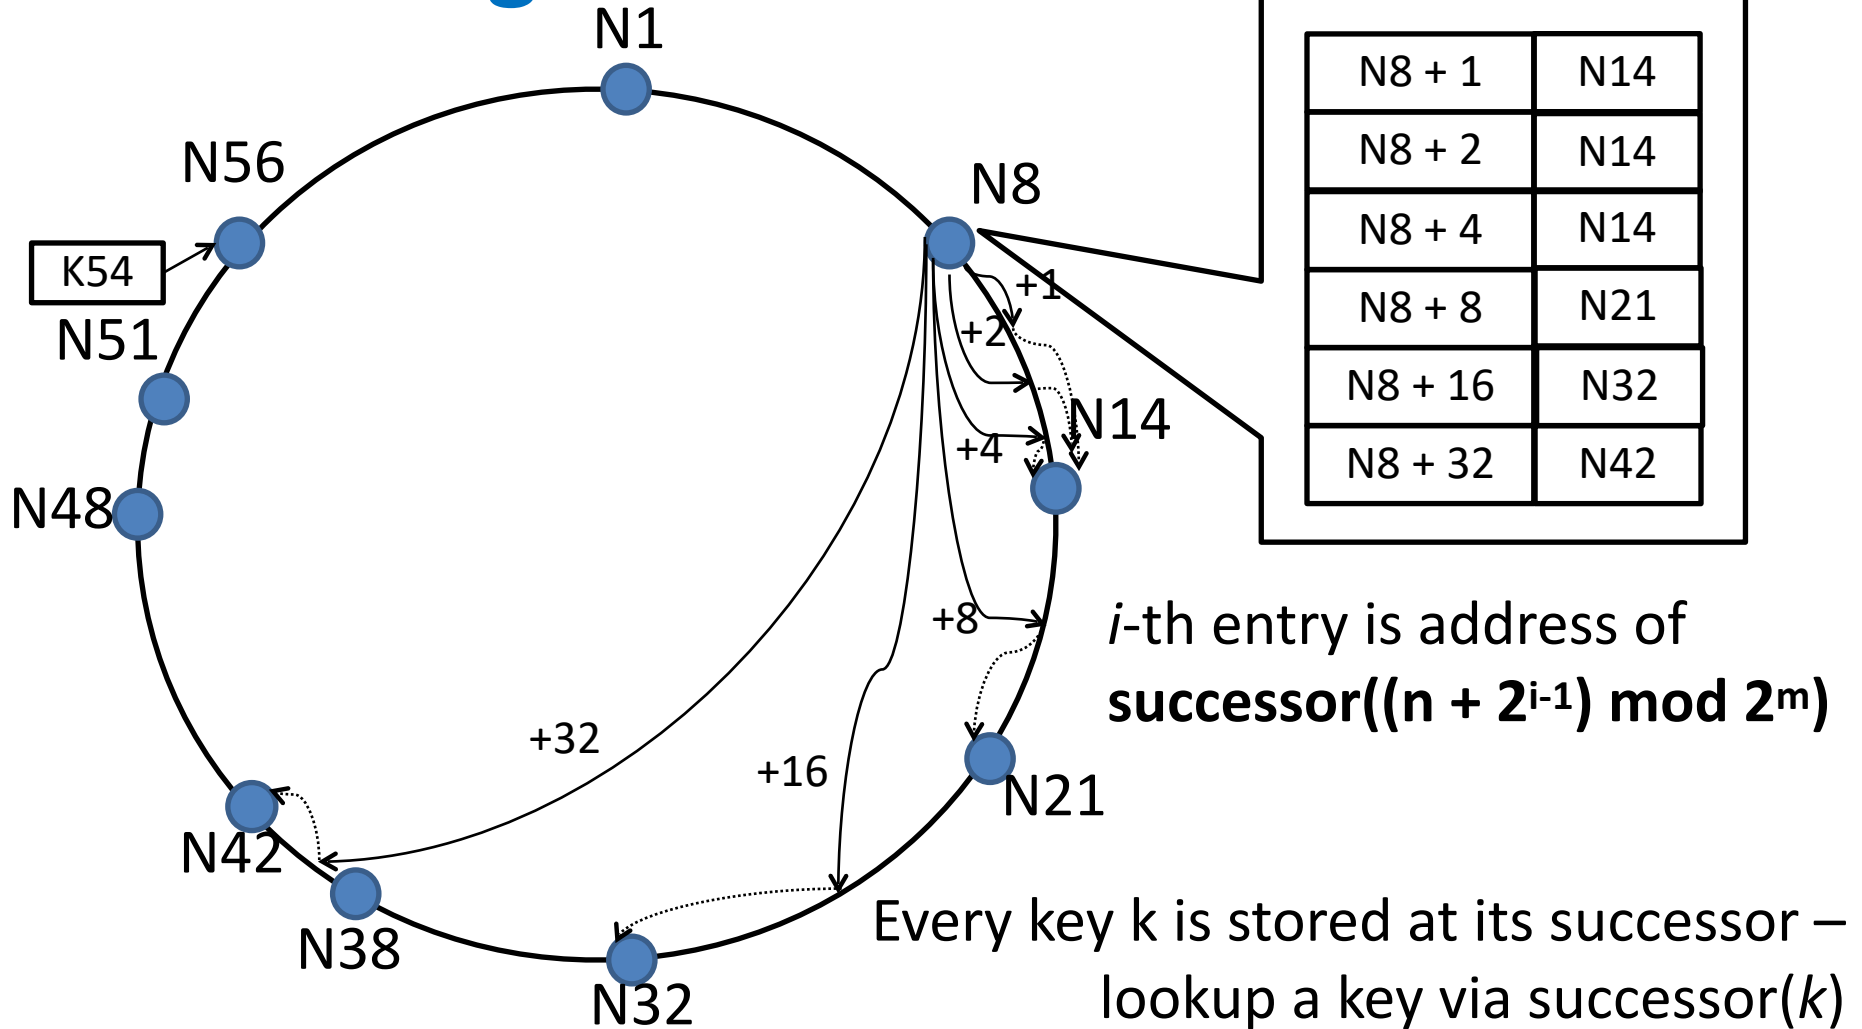
Node Joins & Leaves

- Nodes may disappear from the network (e.g., failure, departure)
- Each node records a whole **segment of the circle** adjacent to it, i.e., **r nodes preceding** and **following** it
- With high probability a node is able to correctly locate its successor or predecessor (even under high node churn)
- When a new node joins or leaves the network, responsibility for $O(K/N)$ keys changes hands

Searching in Chord

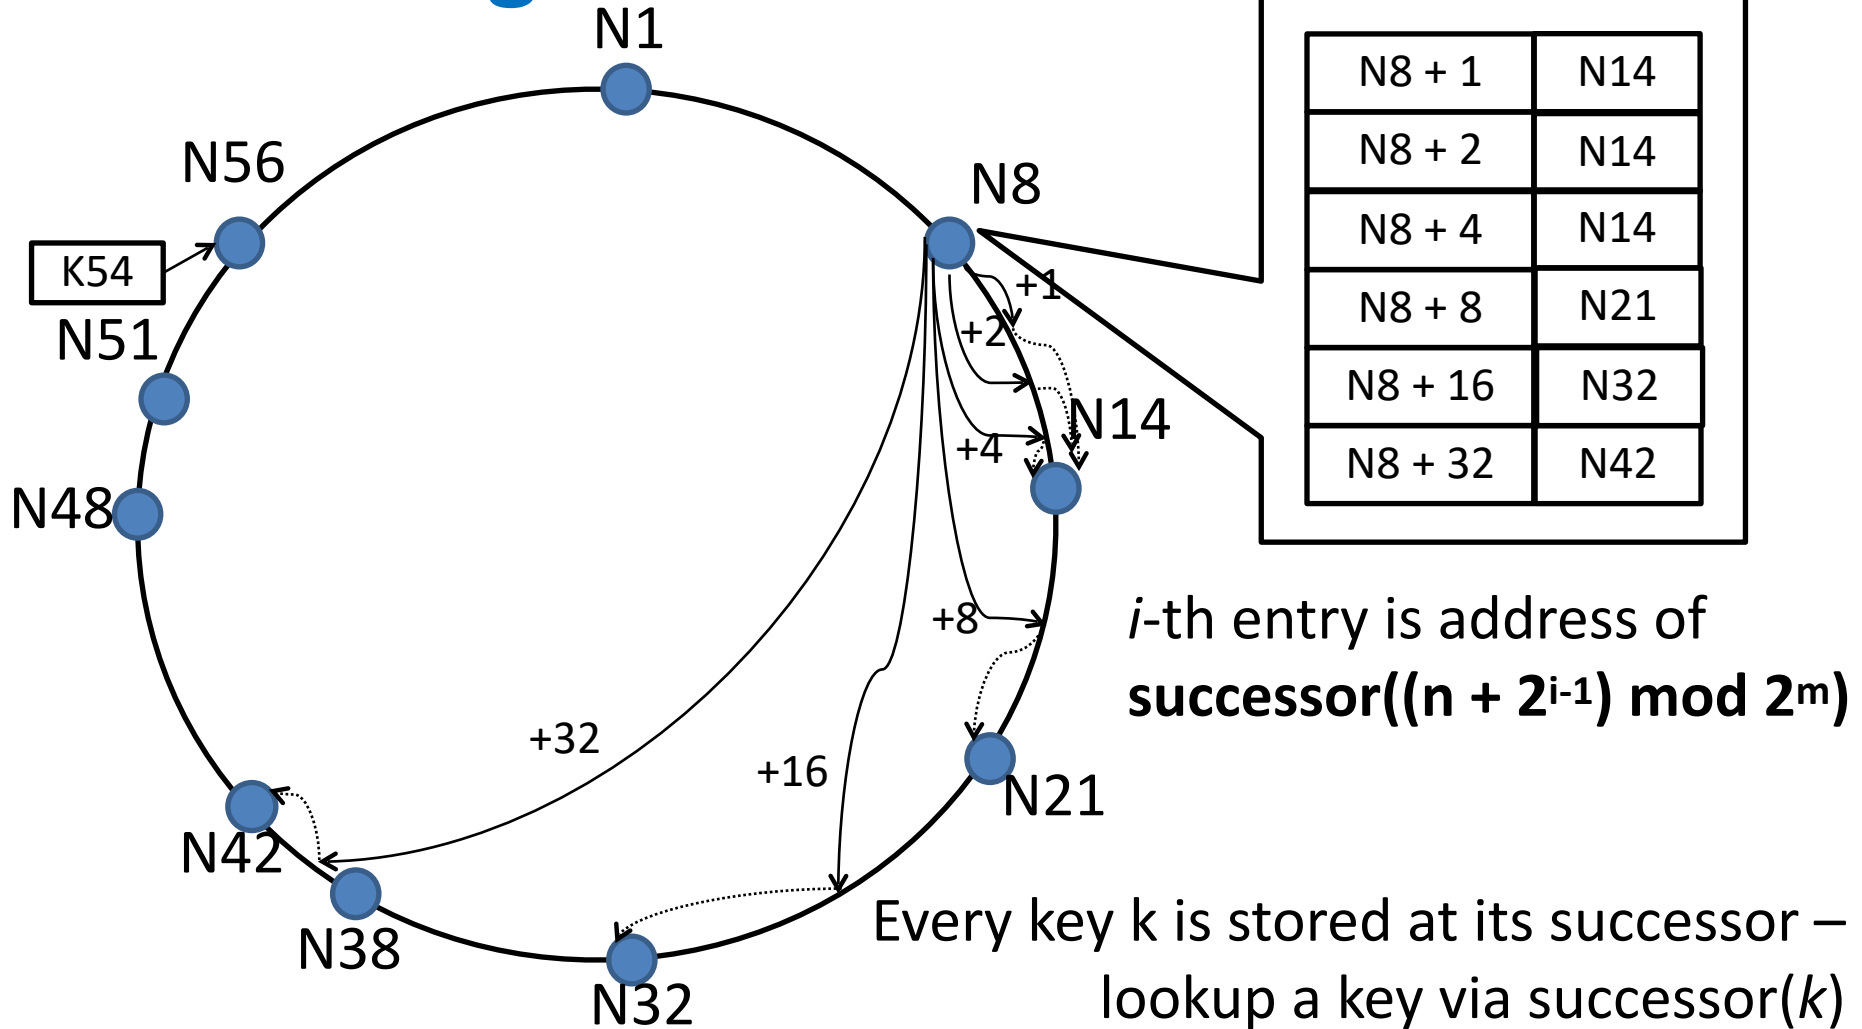
- With knowledge of a single successor, a linear search through network could locate key (naïve search)
- Any given message may potentially have to be relayed through most of the network, i.e., cost is $O(n)$
- Faster search method requires each node to keep a "***finger table***" (FT) containing up to ***m*** entries
 - *i*-th entry in FT of node ***n*** contains the address of **successor((*n* + 2^{i-1}) mod 2^m)**
 - number of nodes that must be contacted to find a successor in an ***n***-node network is **$O(\log n)$**

Chord Finger Table



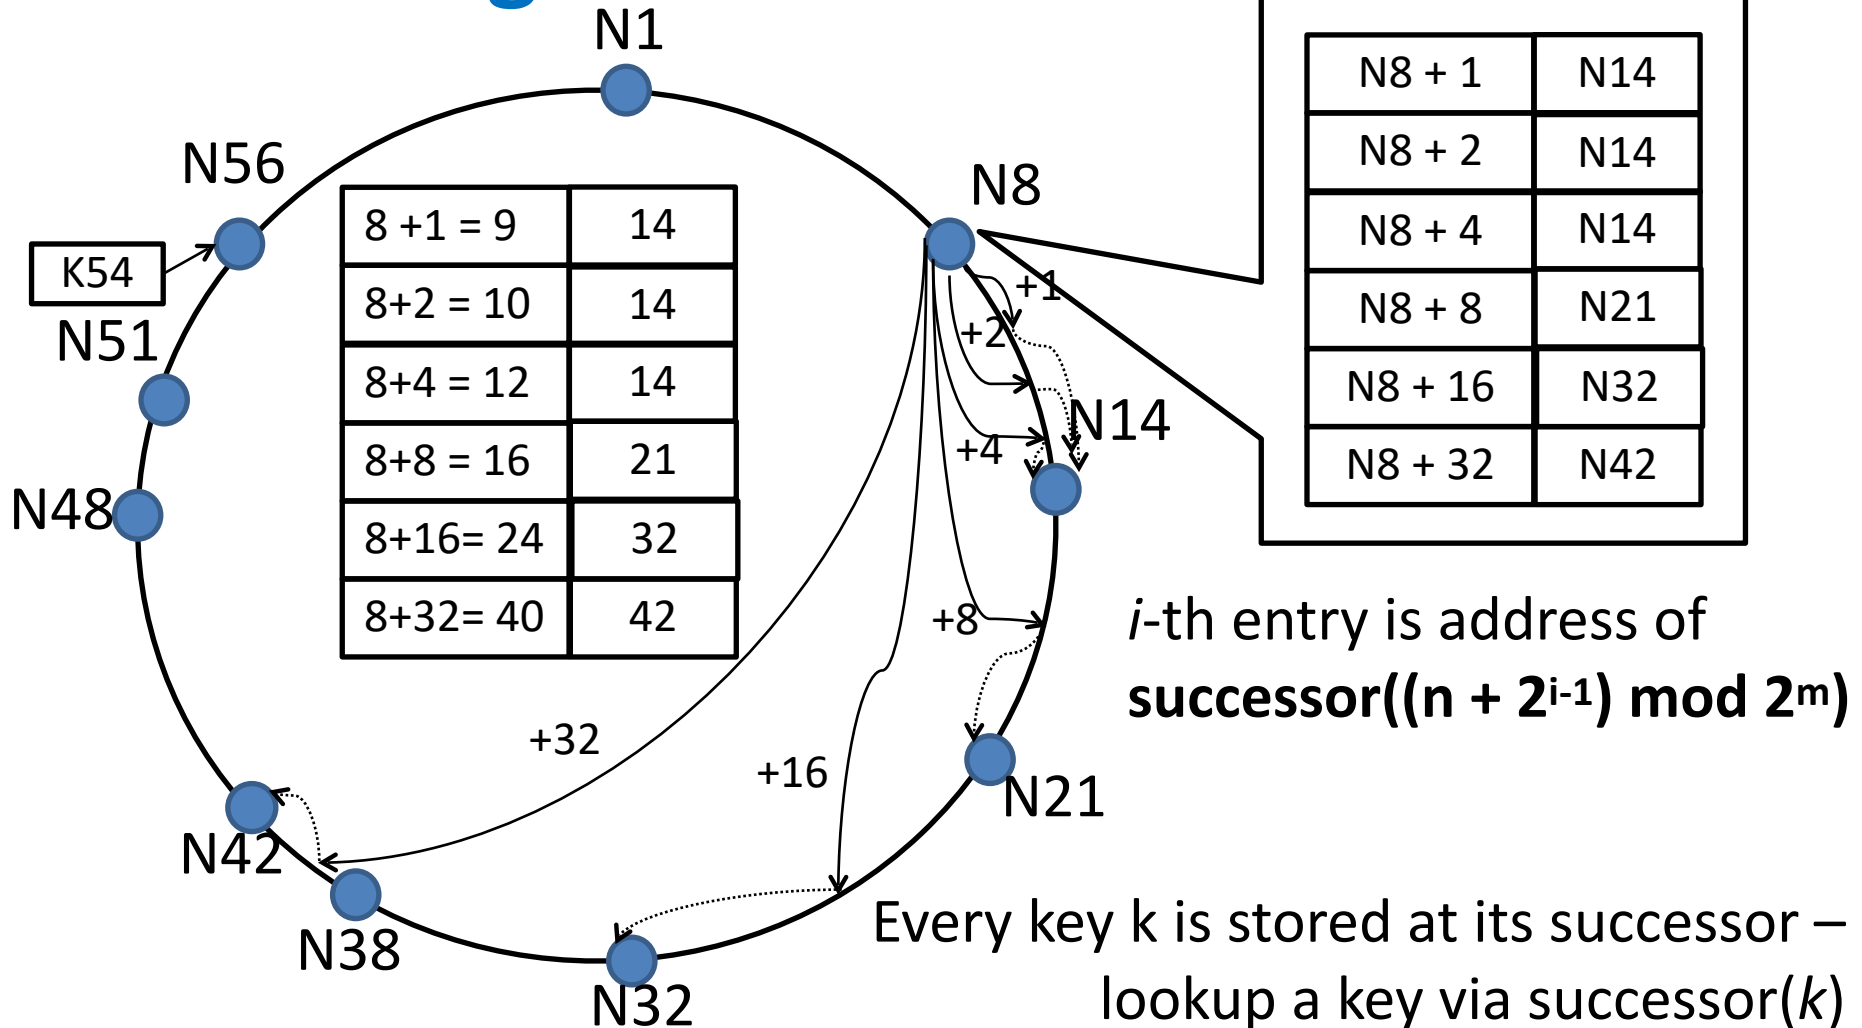
successor(...) is the node on the circle associated with the input argument – whether it is a *key* or a *node identifier*

Chord Finger Table



successor(...) is the node on the circle associated with the input argument – whether it is a *key* or a *node identifier*


Chord Finger Table



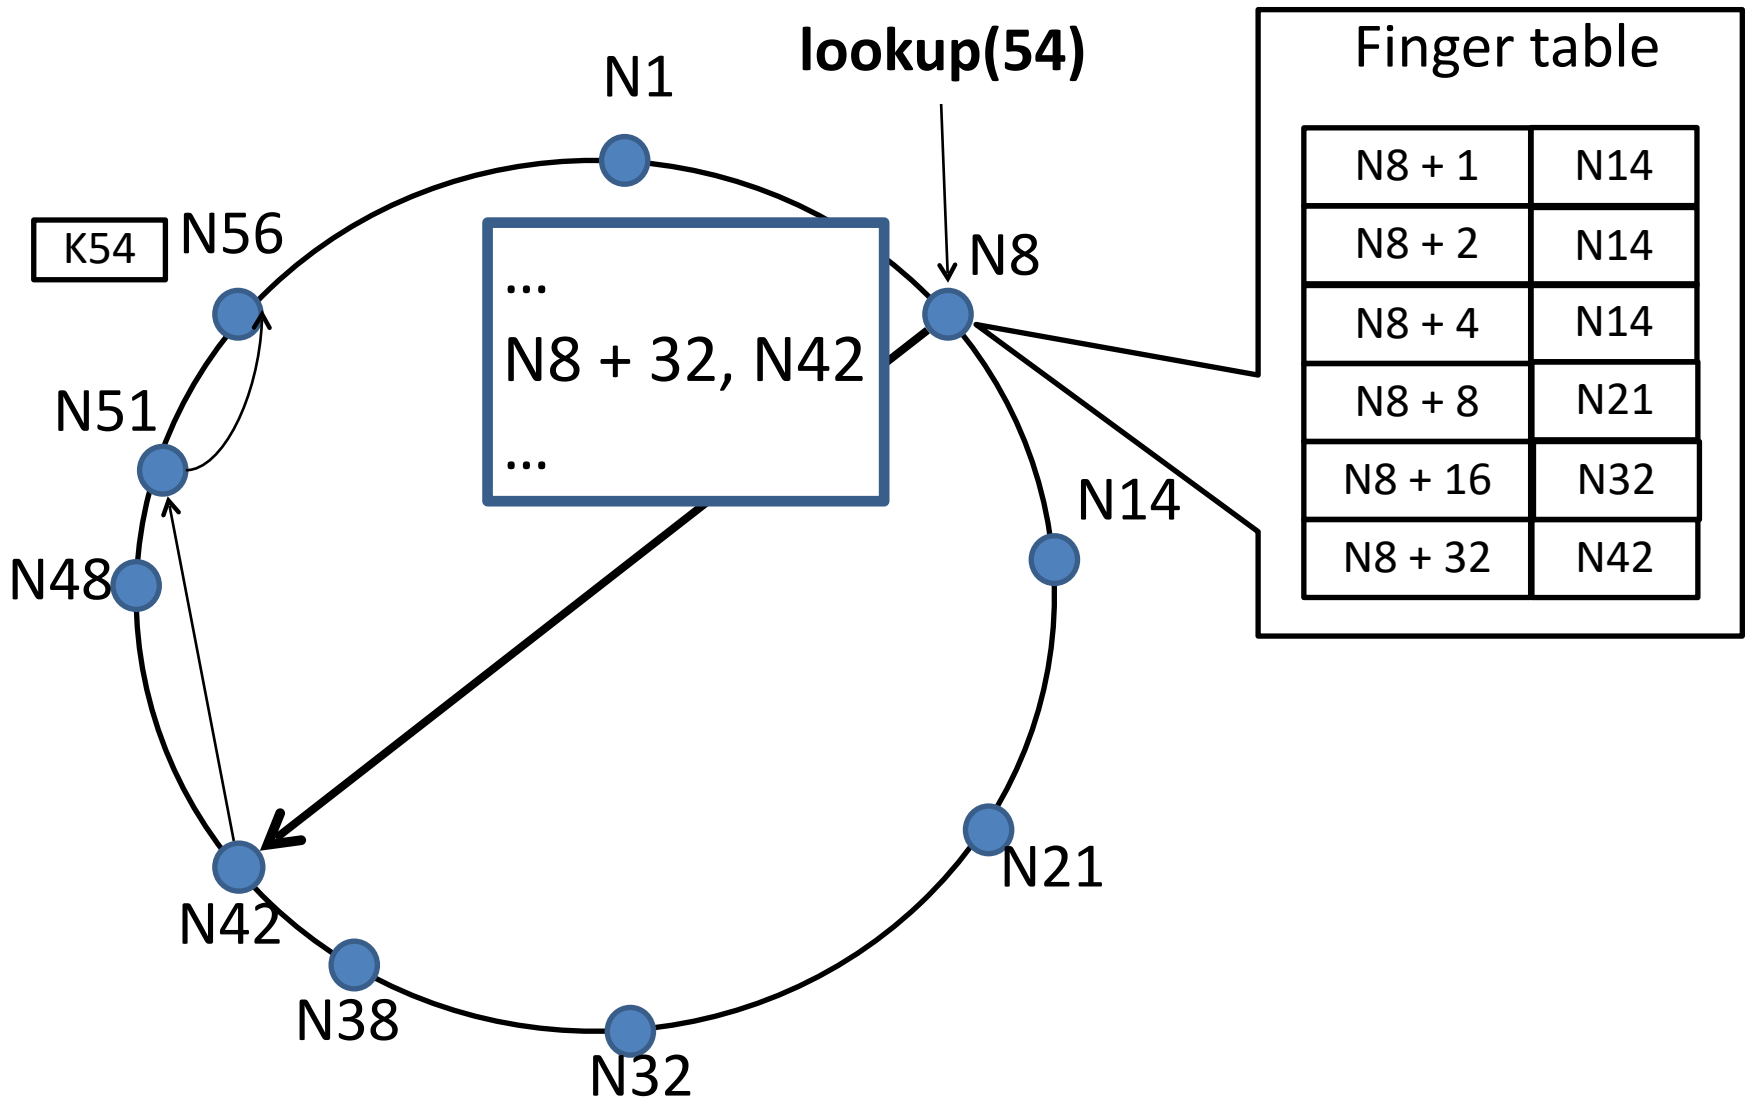
successor(...) is the node on the circle associated with the input argument – whether it is a *key* or a *node identifier*

Chord Key Location

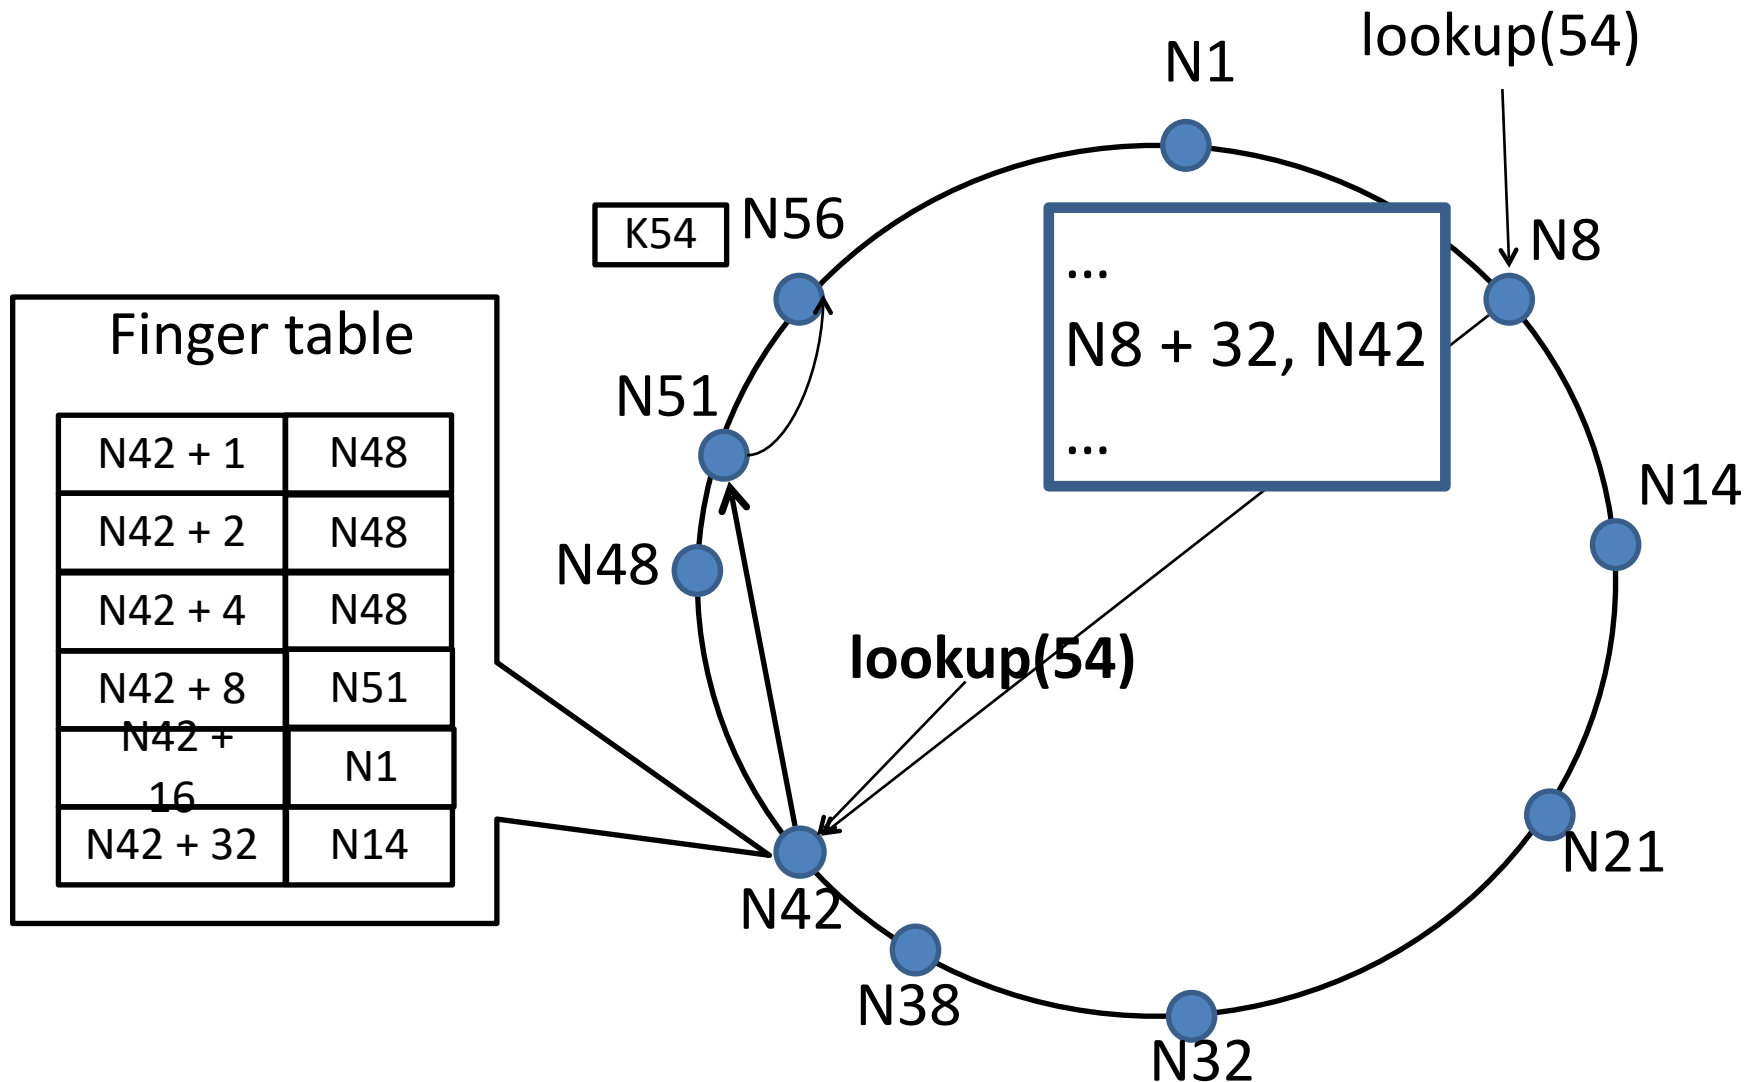
- Lookup in finger table the farthest node that precedes key – closest successor of key in FT
- Query homes in on target in $O(\log n)$ hops
- Each hop at least halves distance to destination



Lookup Example



Lookup Example (cont.)




Lookup Latencies

- While $O(\log n)$ is better than $O(n)$, it can still take considerable amount of time to find the target
- For example, $\log(1,000,000)$ hops which may be distributed anywhere
- Results in potentially high response latencies

Network locality

- Nodes close on ring can be far away in network



* Figure from <http://project-iris.net/dht-toronto-03.ppt>

Peer-to-Peer Systems




CAN Pastry DHT

CAN: Content Addressable Network

- Design is based on virtual multi-dimensional Cartesian coordinate space to organize overlay
- Nodes are mapped into space (coordinates at edges wrap around)
- Address space is independent of physical location and physical connectivity of nodes
- Points in the space are identified with coordinates
- General model is an n -dimensional torus that uses dimensions for routing


CAN: Content Addressable Network

- Design is based on **virtual multi-dimensional Cartesian coordinate space** to organize overlay
- Nodes are mapped into space (**coordinates at edges wrap around**)
- Address space is **independent of physical location** and **physical connectivity** of nodes
- **Points** in the space are **identified with coordinates**
- General model is an n -dimensional torus that uses dimensions for routing



Example 2-d space with 5 nodes



- Entire coordinate space is **dynamically partitioned** among all nodes in system
- Each node owns a distinct **zone** in space
- Each key hashes to a **point** in space



* All CAN figures from "A Scalable Content-Addressable Network", S. Ratnasamy et al., In Proceedings of ACM SIGCOMM 2001.


Coordinates Wrap Around

In all dimensions (only x-axis shown)



CAN Routing

- Put(key, data), get(key)
- Greedily forward message to **neighbor closest to destination** in Cartesian coordinate space
- Nodes maintain a **routing table** that holds IP address and zone of its neighbours
1's coordinate neighbor set = {2,3,4,5}



Sample routing path from node 1 to point (x,y)

CAN Routing

- Many possible routing paths exist between two points in space
- If a neighbour on a path crashes, simply pick the next best available (node) path


Average Path Length

- **d -dimensional space**, partitioned into **n equal-sized zones**, **average routing path length** is: **$d/4 * n^{1/d}$**
- Each node maintains $2d$ neighbours
- Grow number of nodes, without affecting per node state
- Grow number of nodes, increases path length by $O(n^{1/d})$
- 2-dimensional space: $1/2 * n^{1/2}$ (average routing path)
- 3-dimensional space: $3/4 * n^{1/3}$ (average routing path)

Node Joining a CAN

- Find a node already in overlay network
- Identify a zone that can be split
 - Pick random point
 - Route join request to node managing the point's zone
 - Initiate split of zone at that node
- Update routing tables of nodes neighbouring newly split zone
- If refused, try with a new random point


Zone Splitting Upon Node Joining



1's coordinate
neighbor
set = {2,3,4,5}

Join request
Node 7

Zone Splitting Upon Node Joining




1's coordinate
neighbor
set = {2,3,4,5}

Join request
Node 7

- Pick random
point (A, B)

Zone Splitting Upon Node Joining




1's coordinate
neighbor
set = {2,3,4,5}

Join request
Node 7

- Pick random point (A, B)
- Route join request of 7 to 1

Zone Splitting Upon Node Joining




1's coordinate
neighbor
set = {2,3,4,5}

Join request
Node 7

- Pick random point (A, B)
- Route join request of 7 to 1
- Initiate zone split

Zone Splitting Upon Node Joining




1's coordinate
neighbor
set = {2,3,4,5}

Join request
Node 7

- Pick random point (A, B)
- Route join request of 7 to 1
- Initiate zone split

Zone Splitting Upon Node Joining



1's coordinate
neighbor
set = {2,3,4,7}

7's coordinate
neighbor
set = {1,2,4,5}

Update routing
tables of nodes,
transfer state,
i.e., (k, v)-pairs
(not shown)

Node Join Properties

- Only $O(d)$ nodes are effected when a node joins/leaves CAN (a node has $2d$ neighbours)
- Independent of n , number of nodes in CAN

Pastry (2001)

- Ring-based partitioning like Chord
- Each peer discovers and exchanges state information: List of leaf nodes, neighborhood list, routing table
- **Leaf node list** are $L/2$ closest peers by Node ID in each direction around the circle
 - Lookups first search the leaf node list
- **Neighborhood list** are M closest peers in terms of routing metric (e.g. ping delay)
 - Good candidates for routing table

Routing Table

- 6 digits, base 4: table of 6x4 entries
- Row i contains nodes which share $i-1$ -th long prefix
- Populate cells with neighbors if possible
- Column indicates i -th digit
- Lookup in RT finds a node with a longer prefix

<u>Node 103220</u>	0	1	2	3
1	031120	103220	201303	312201
2	103220	110003	120132	132012
3	100221	101203	102303	103220
4		103112	103220	103302
5		103210	103220	
6	103220			

Example: lookup(102332) -> 102303

DHT Routing Summary

- Chord
 - Finger table routing
 - Each hop at least halves distance (in identifier circle) to destination
- Pastry
 - Proximity-based Routing
- CAN
 - Neighbour routing
 - Forward to neighbor closest (in Cartesian coordinate space) to destination

Conclusions on P2P

- Hugely popular area of research 2000-2010
- Large-scale companies (Amazon & Google et al.) prefer self-managed cloud infrastructures
- P2P principles, techniques and abstractions are used by large-scale systems (e.g., **DHTs**)
- Active applications: **BitTorrent**, **Bitcoin** *et al.*
- Peer-assisted, hybrid systems were popular: **Skype**, **Spotify**, etc.